

CS 480

Introduction to Artificial Intelligence

September 2nd, 2021

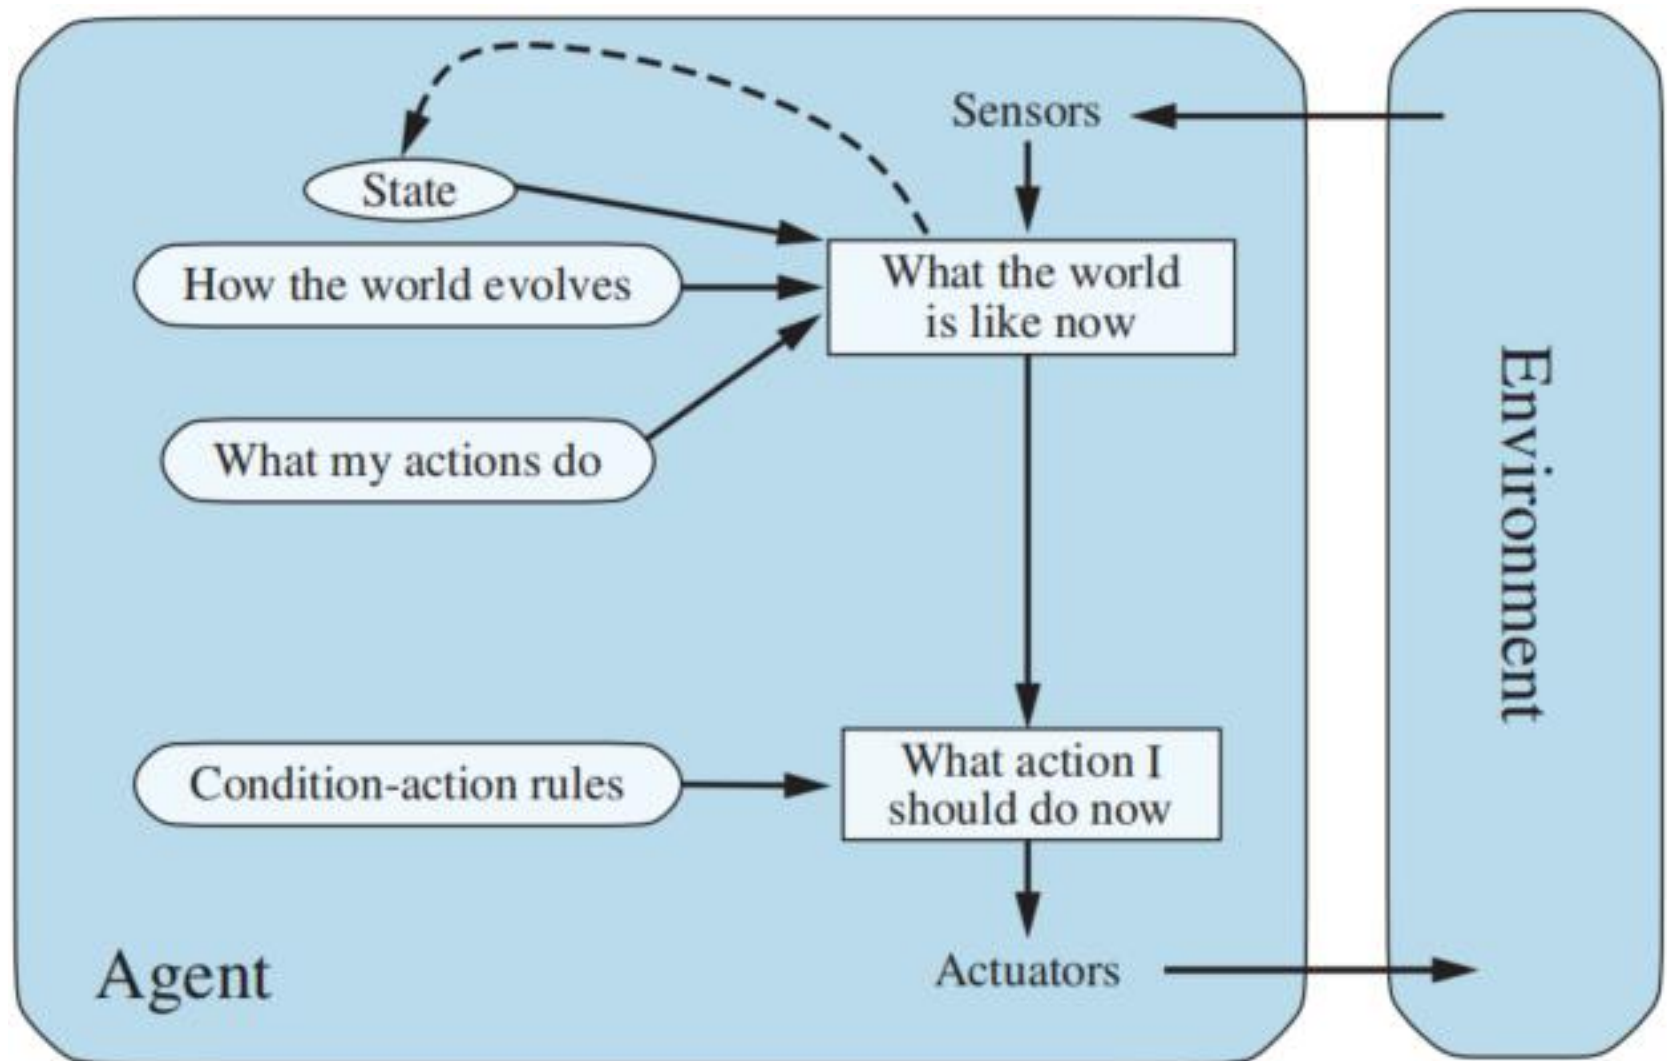
Announcements / Reminders

- **Contribute to the discussion on Blackboard, please**
- **Please follow the Week 02 To Do List instructions**

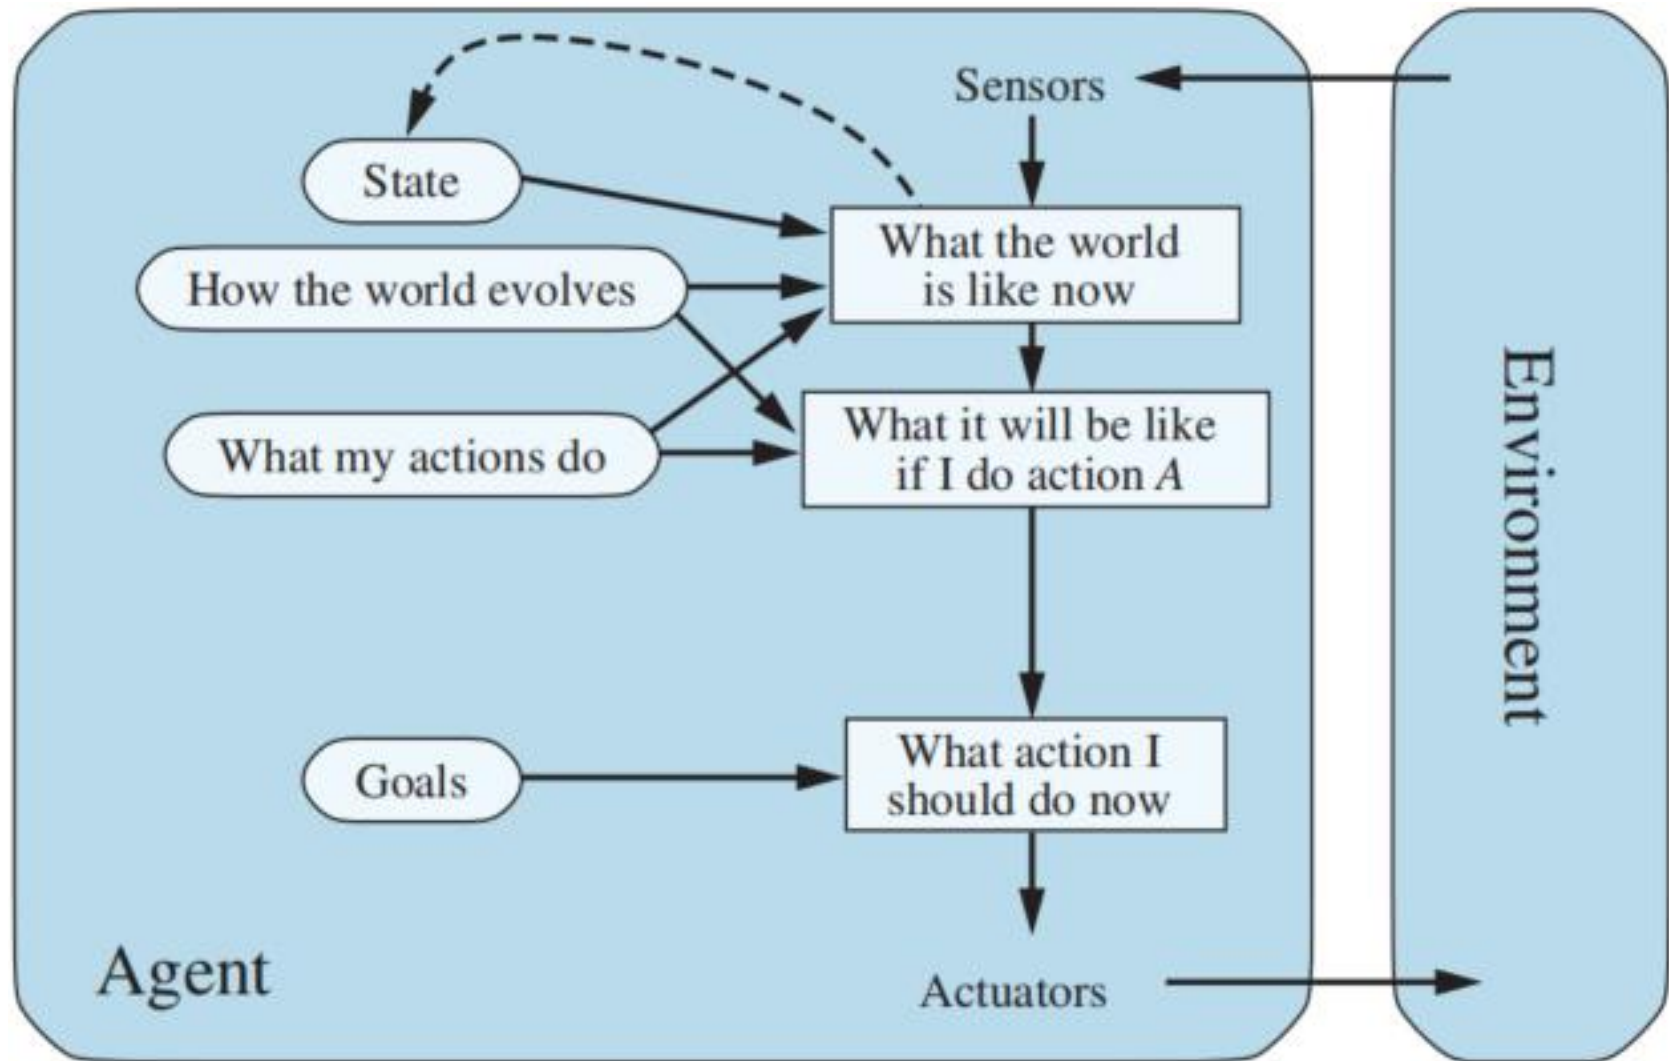
Plan for Today

- **Intelligent Agents**
- **Problem Solving: Searching**


Model-based Reflex Agent



Model-based Goal-based Agent



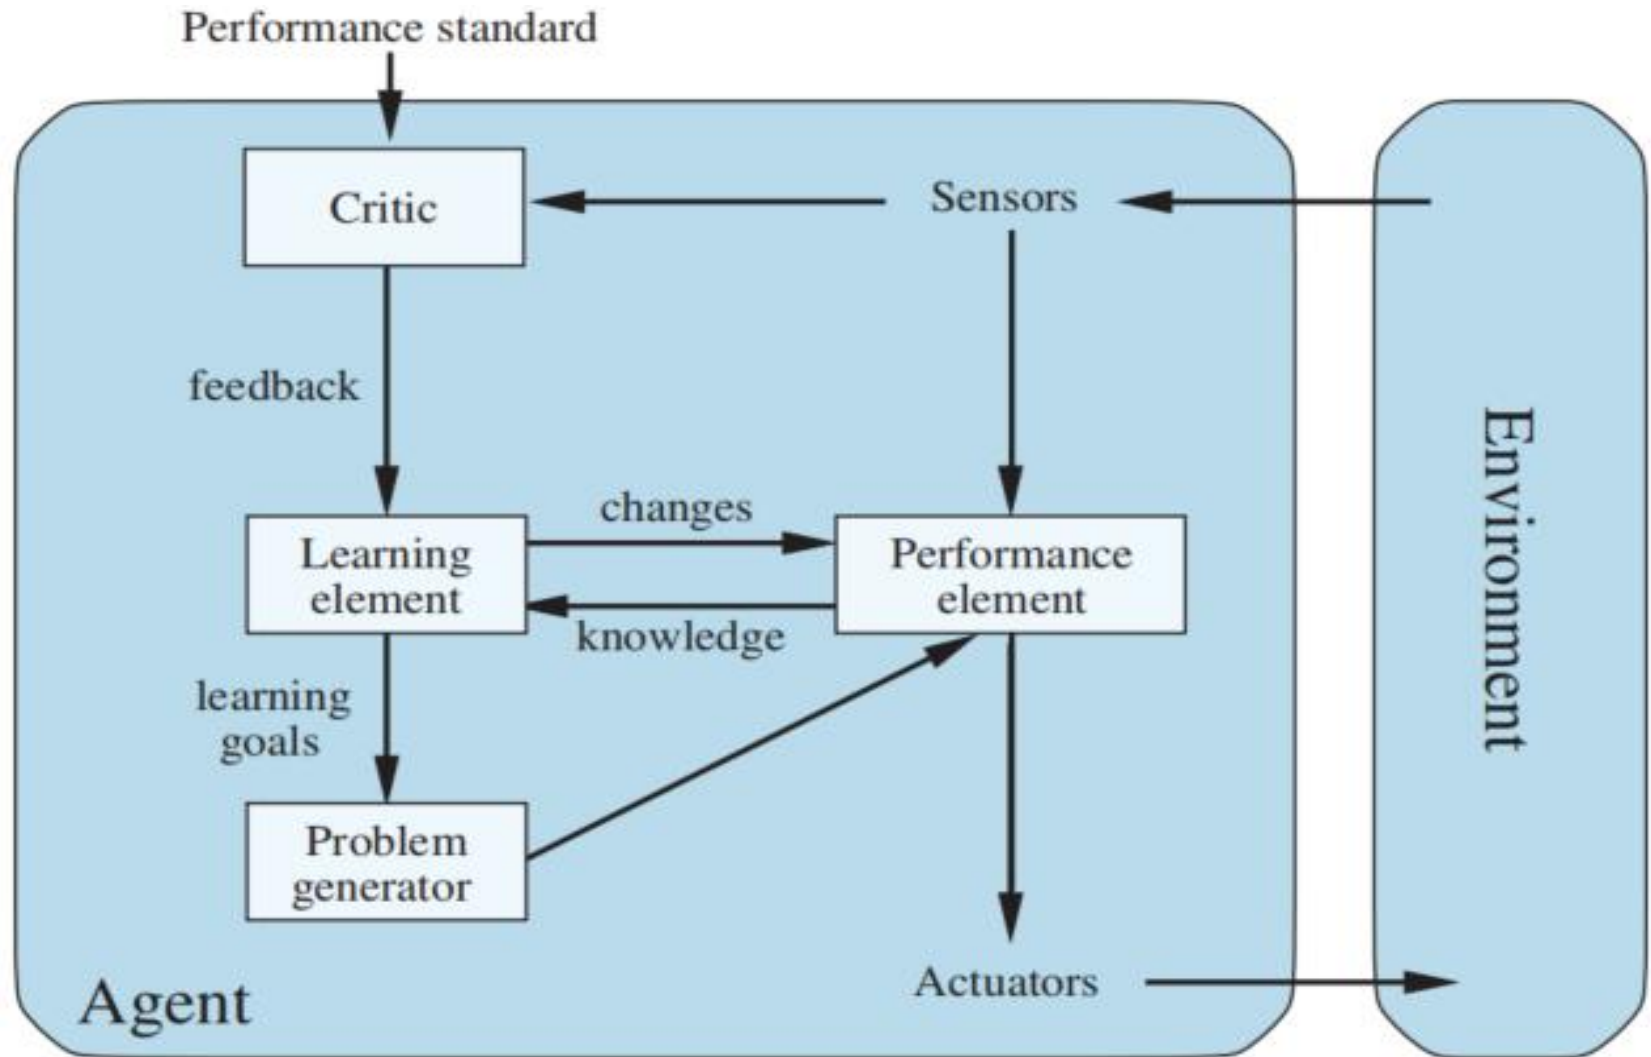
Model-based Utility-based Agent



Typical Agent Architectures

- Simple reflex agent: uses condition-action rules
- **Model-based** reflex agent: keeps track of the unobserved parts of the environment by maintaining internal state:
 - “how the world works”: state transition model
 - how percepts and environment is related: sensor model
- Goal-based reflex agent: maintains the **model of the world** and goals to select decisions (that lead to goal)
- Utility-based reflex agent: maintains the **model of the world** and utility function to select PREFERRED decisions (that lead to the best expected utility: $\text{avg}(\text{EU} * p)$)

Learning Agent



Designing the Agent for the Task


**Analyze the
Problem / Task
(PEAS)**

**Select Agent
Architecture**

**Select Internal
Representations**


**Apply
Corresponding
Algorithms**

State and Transition Representations




- **Atomic:** state representation has NO internal structure
- **Factored:** state representation includes fixed attributes (which can have values)
- **Structured:** state representation includes objects and their relationships


State and Transition Representations



(a) Atomic




(b) Factored




(c) Structured

Complexity, level of detail, expresiveness, more difficult to process


Representations and Algorithms



(a) Atomic



(b) Factored



(c) Structured

- Searching
- Hidden Markov models
- Markov decision process
- Finite state machines

- Constraint satisfaction algorithms
- Propositional logic
- Planning
- Bayesian algorithms
- Some machine learning algorithms

- Relational database algorithms
- First-order logic
- First-order probability models
- Natural language understanding (some)

Finite State Machine: A Turnstile


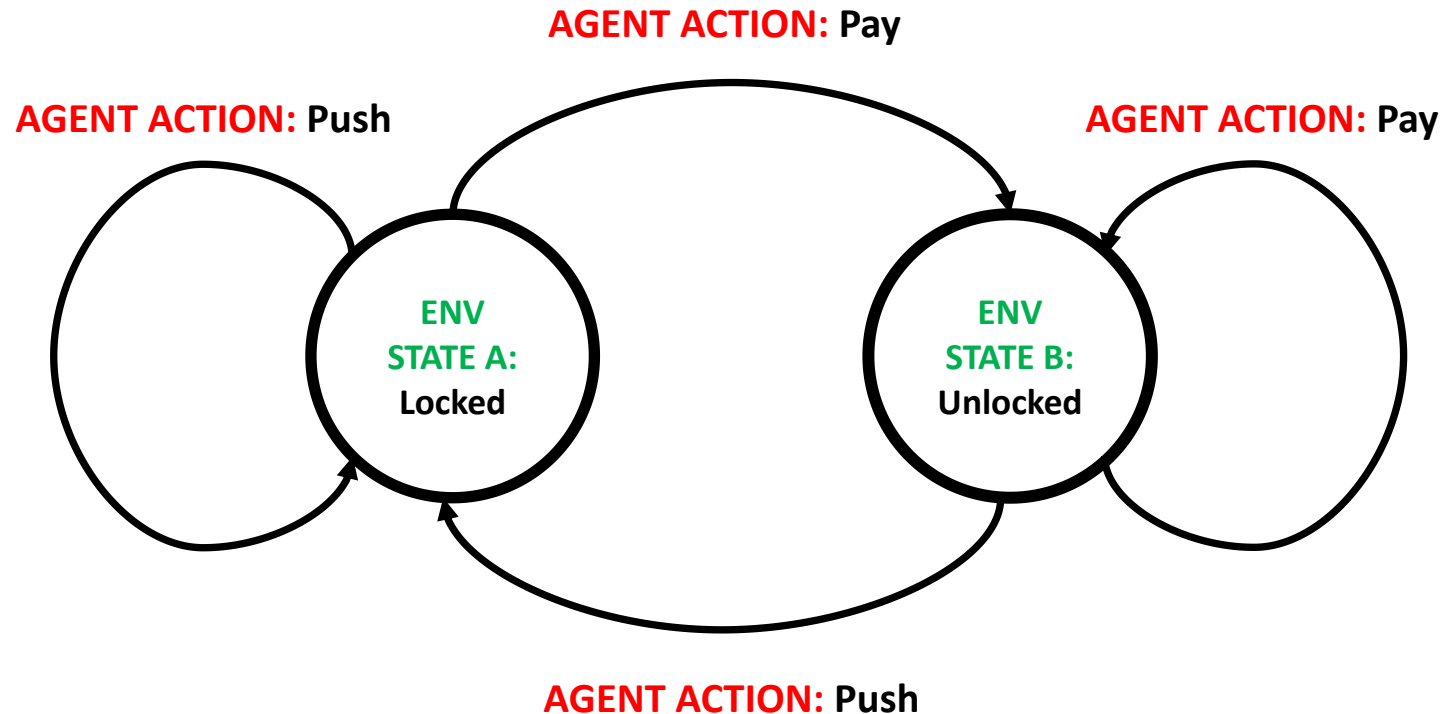
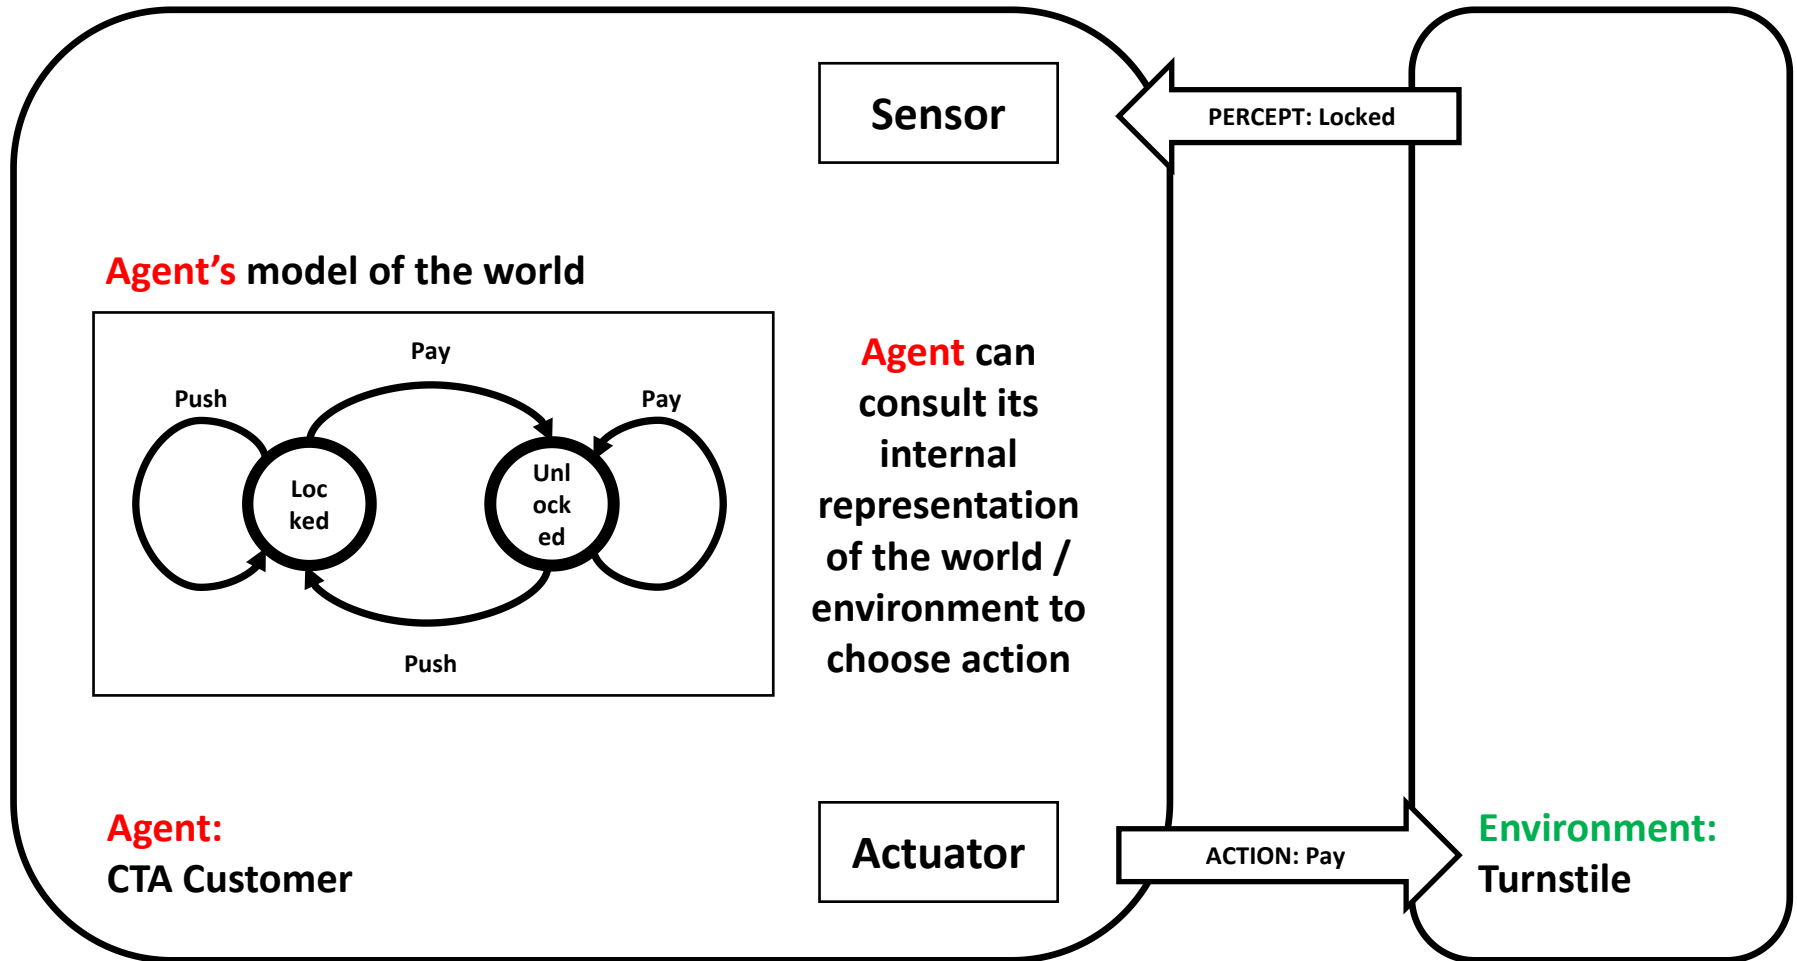


Image source: Wikipedia

Finite State Machine: A Turnstile

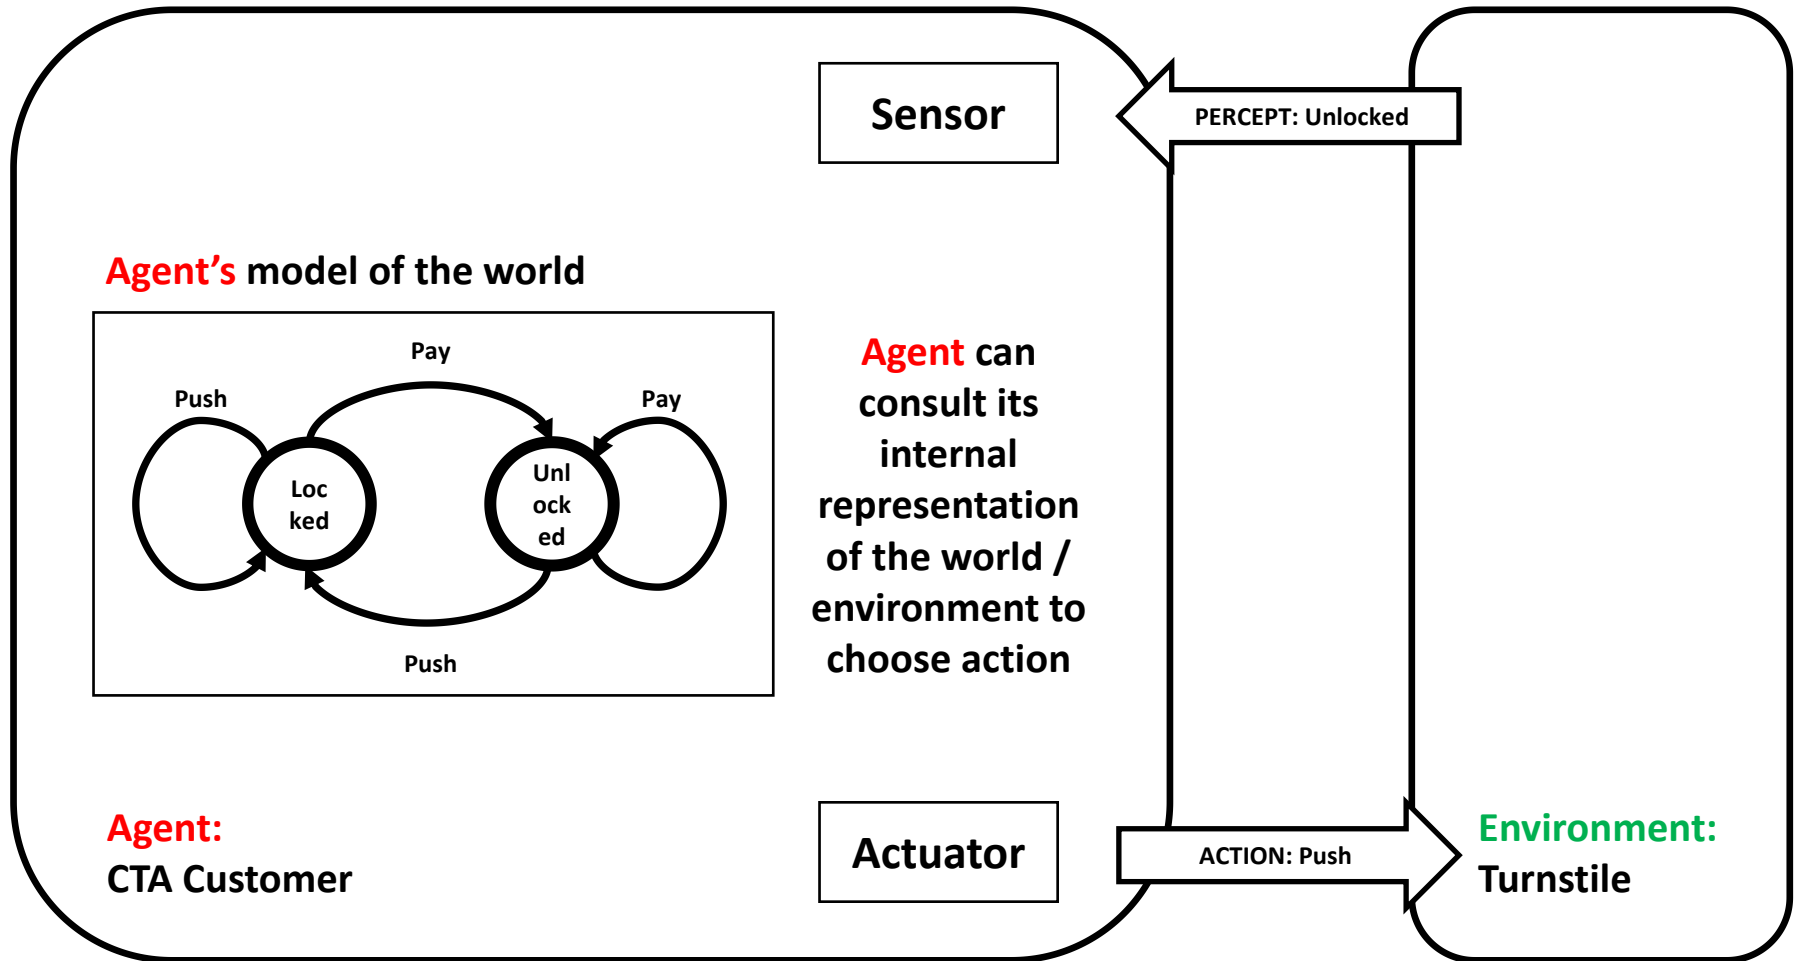


Model-based Reflex Agent Example




Note: This problem could be easily solved with a simple (without internal model) reflex agent.

Model-based Reflex Agent Example




Note: This problem could be easily solved with a simple (without internal model) reflex agent.


Representations: Examples






(a) Atomic



(b) Factored



(c) Structured



Designing the Agent for the Task

**Analyze the
Problem / Task
(PEAS)**

**Select Agent
Architecture**

**Select Internal
Representations**

**Apply
Corresponding
Algorithms**

BTW: How Would you Program it All?

Problem-Solving / Planning Agent

- **Context / Problem:**

- correct action is NOT immediately obvious
- a plan (a sequence of actions leading to a goal) may be necessary

- **Solution / Agent:**

- come up with a computational process that will search for that plan

- **Planning Agent:**

- uses factored or structured representations of states
- uses searching algorithms

Planning: Environment Assumptions

Works with a “Simple Environment”:

- Fully observable
- Single agent (for now -> it can be multiagent)
- Deterministic
- Static
- Episodic
- Discrete
- Known to the agent

Problem-Solving Process

- **Goal formulation:**
 - adopt a goal (think: desirable state)
 - a concrete goal should help you reduce the amount of searching
- **Problem formulation:**
 - an **abstract** representation of states and actions
- **Search:**
 - search for solutions within the **abstract** world model
- **Execute actions in the solution**

Planning: Environment Assumptions

Works with a “Simple Environment”:

- Fully observable
- Single agent (for now -> it can be multiagent)
- Deterministic
- Static
- Episodic
- Discrete
- Known to the agent

Important and helpful:

Such assumptions **GUARANTEE** a **FIXED** sequence of actions as a solution


What does it mean?

You can execute the “plan” without worrying about incoming percepts (open-loop control)

Defining Search Problem

- Define a set of possible states: **State Space**
- Specify **Initial State**
- Specify **Goal State(s)** (there can be multiple)
- Define a FINITE set of possible **Actions** for EACH state in the State Space
- Come up with a **Transition Model** which describes what each action does
- Specify the **Action Cost Function**: a function that gives the cost of applying action a in state s

Sample Problem: Dracula's Roadtrip




Problem: Get from Arad to Bucharest efficiently (for example: quickly or cheaply).


Search Problem: Dracula's Roadtrip

- State Space: **a map of Romania**
- Initial State: **Arad**
- Goal State: **Bucharest**
- Actions:
 - for example: **ACTIONS(Arad) = {ToSibiu, ToTimisoara, ToZerind}**
- Transition Model:
 - for example: **RESULT(Arad, ToZerind) = Zerind**
- Action Cost Function [**ActionCost(S_{current} , a, S_{next})**]
 - for example: **ActionCost(Arad, ToSibiu, Sibiu) = 140**


Sample Problem: Dracula's Roadtrip



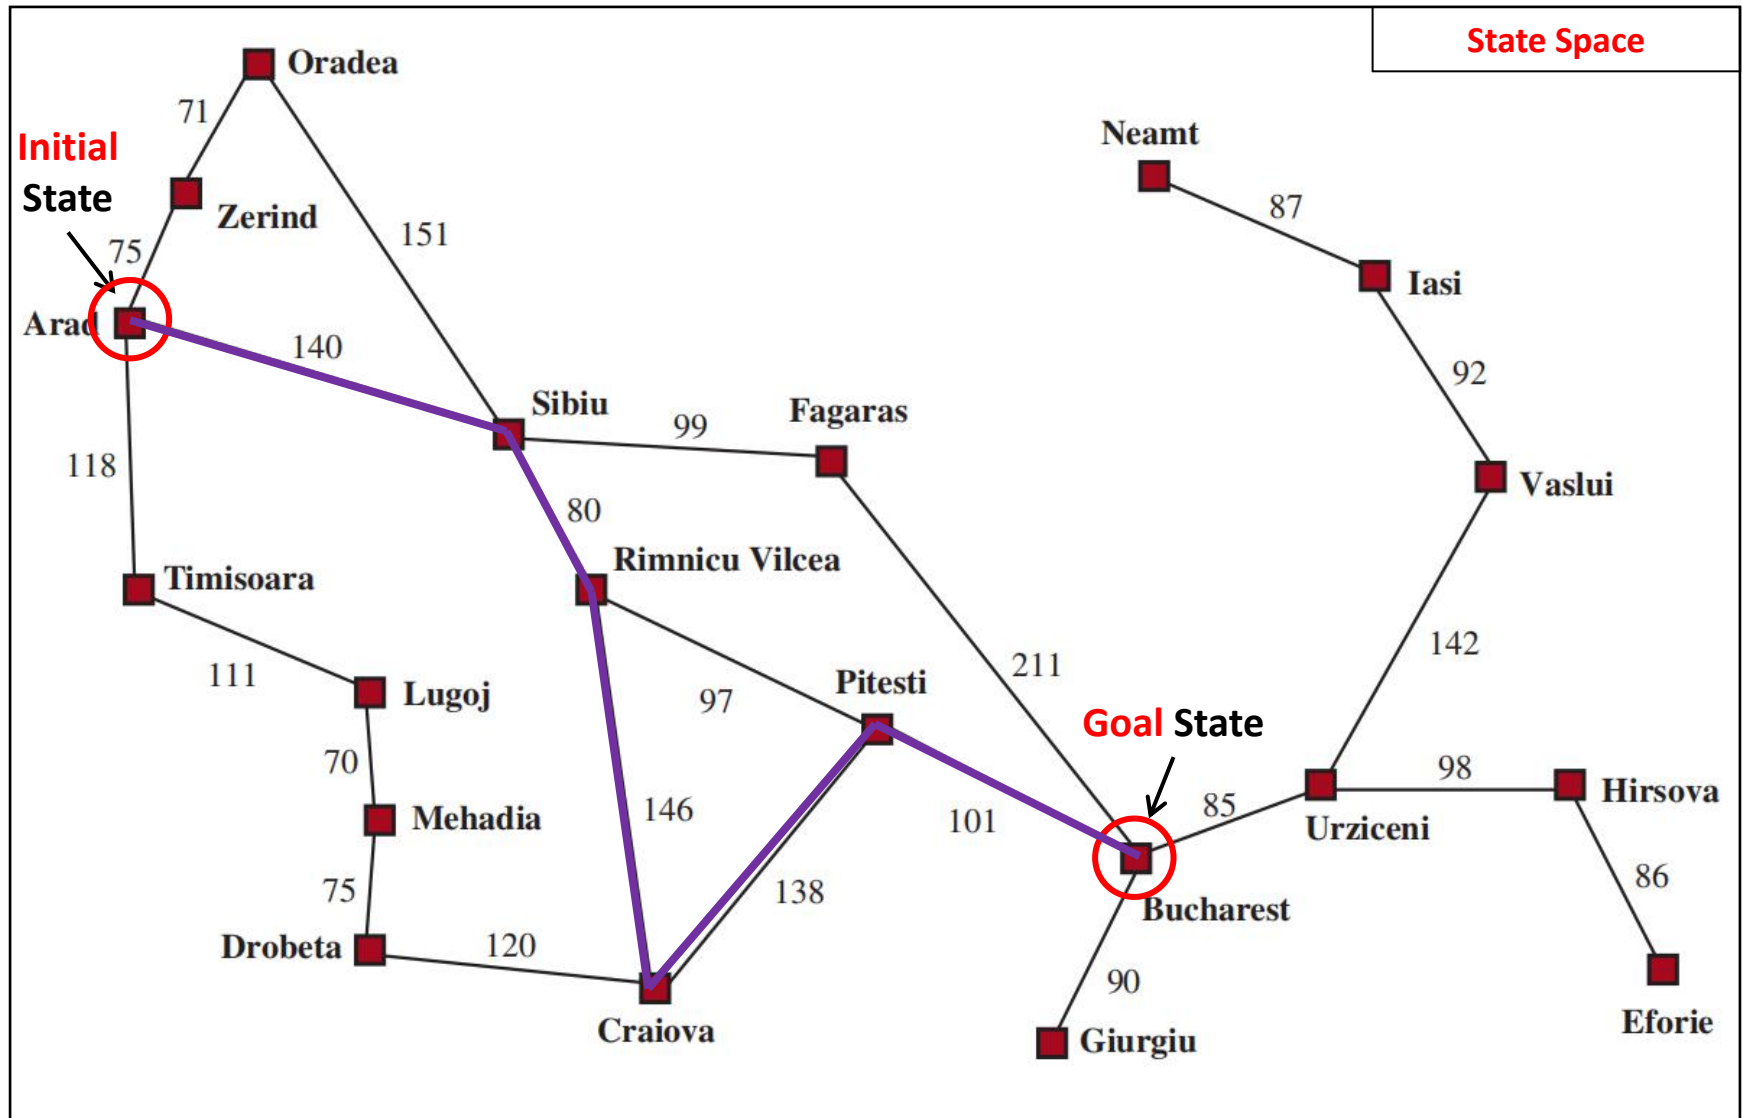
Dracula's Roadtrip: Potential Solution




Dracula's Roadtrip: Potential Solution




Dracula's Roadtrip: Potential Solution




Tic Tac Toe: (Partial) State Space




Tic Tac Toe: Solution




Initial State




X's move




O's move




X's move




O's move



X's move



O's move




X's move

Solution



A **Solution** is a sequence of actions (a path) between the initial state and the goal state

Chess: (First Move) State Space


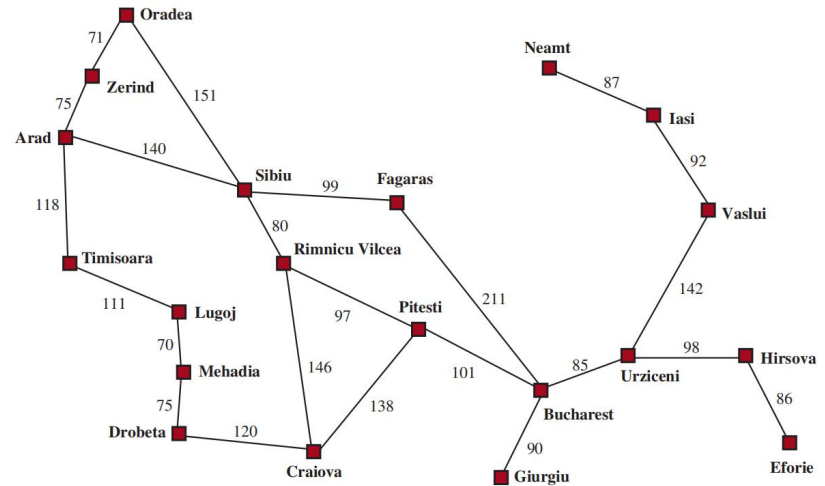

Initial
State




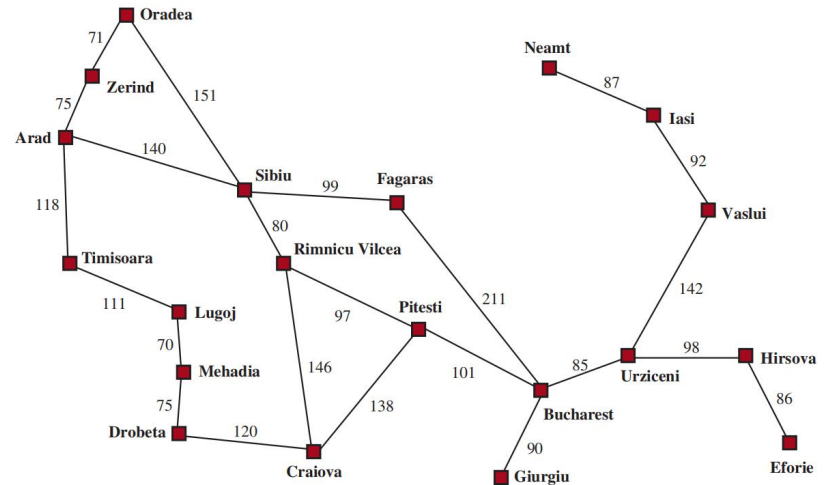

20 Possible **legal** first moves:
16 pawn moves
4 knight moves



Does This Look Familiar?



Does This Look Familiar?



They are all graphs (some will be trees) with **states as nodes** and **actions as links / edges**.

Selected Searching Algorithms

