

# **CS 480**

## ***Introduction to Artificial Intelligence***

**November 18th, 2021**

# Announcements / Reminders

- **Programming Assignment #02:**
  - due on Tuesday, December 7th, at 11:00 PM CST
- **Written Assignment #03:**
  - due on Wednesday, December 1st, at 11:00 PM CST
- **Final Exam:**
  - Thursday, December 2nd, 2021 (during lecture time)

# Plan for Today

- **Casual Introduction to Reinforcement Learning**

# Main Machine Learning Categories

## Supervised learning

**Supervised learning** is one of the most common techniques in machine learning. It is based on **known relationship(s) and patterns within data** (for example: relationship between inputs and outputs).

Frequently used types: **r e g r e s s i o n**, and **classification**.

## Unsupervised learning

**Unsupervised learning** involves finding underlying patterns within data. Typically used in **clustering** data points (similar customers, etc.)

## Reinforcement learning

Reinforcement learning is inspired by behavioral psychology. It is **based on a rewarding / punishing an algorithm**.

Rewards and punishments are based on algorithm's action within its environment.

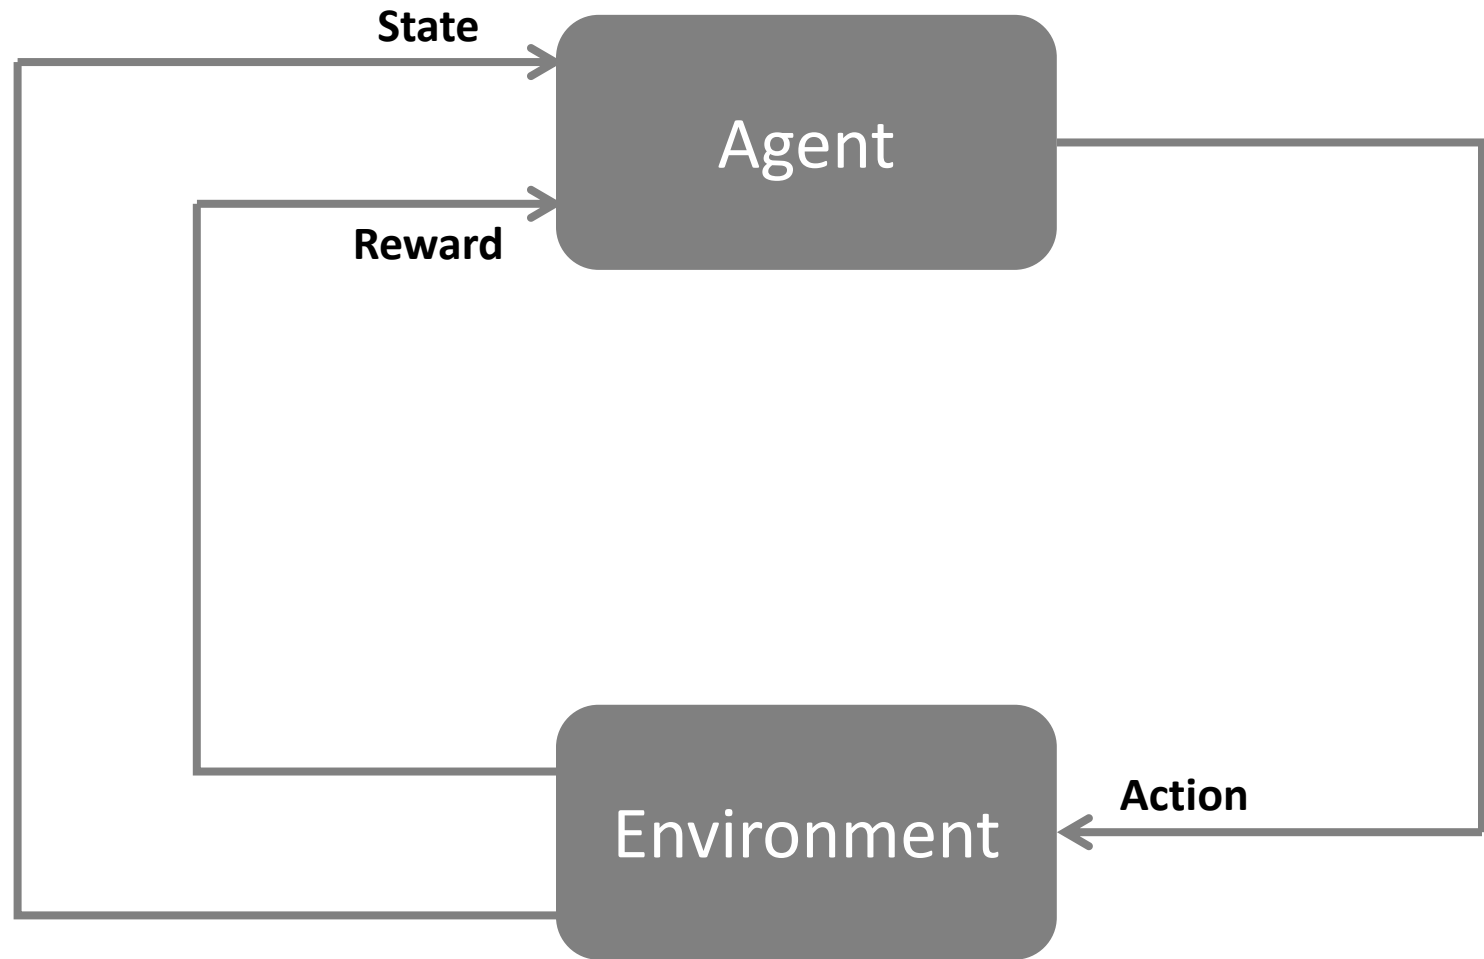
# What is Reinforcement Learning?

## Idea:


Reinforcement learning is inspired by behavioral psychology. It is **based on a rewarding / punishing an algorithm.**

Rewards and punishments are based on algorithm's action within its environment.

# RL: Agents and Environments



# Reinforcement Learning in Action



Solving Rubik's Cube with a Robot Hand

409,438 views • Oct 15, 2019

9.7K

127

SHARE

SAVE

...

Source: <https://www.youtube.com/watch?v=x4O8pojMF0w>



# Reinforcement Learning in Action



Source: <https://www.youtube.com/watch?v=kopoLzvh5jY>



# K-Armed Bandit Problem



# K-Armed Bandit Problem

The K-armed bandit problem is a problem in which a **fixed limited set of resources** must be **allocated between competing (alternative) choices** in a way that **maximizes their expected gain**.

Each choice's **properties are only partially known** at the time of allocation, and **may become better understood as time passes** or by allocating resources to the choice.

# K-Armed Bandit Problem

In the problem, **each machine provides a random reward from a probability distribution specific to that machine, that is not known a-priori.**

The objective of the gambler is to **maximize the sum of rewards earned through a sequence of lever pulls.**

# K-Armed Bandit Problem

Bandit/Arm 1

33 %

**current**  
success (win)  
rate

Bandit/Arm 2

52 %

**current**  
success (win)  
rate

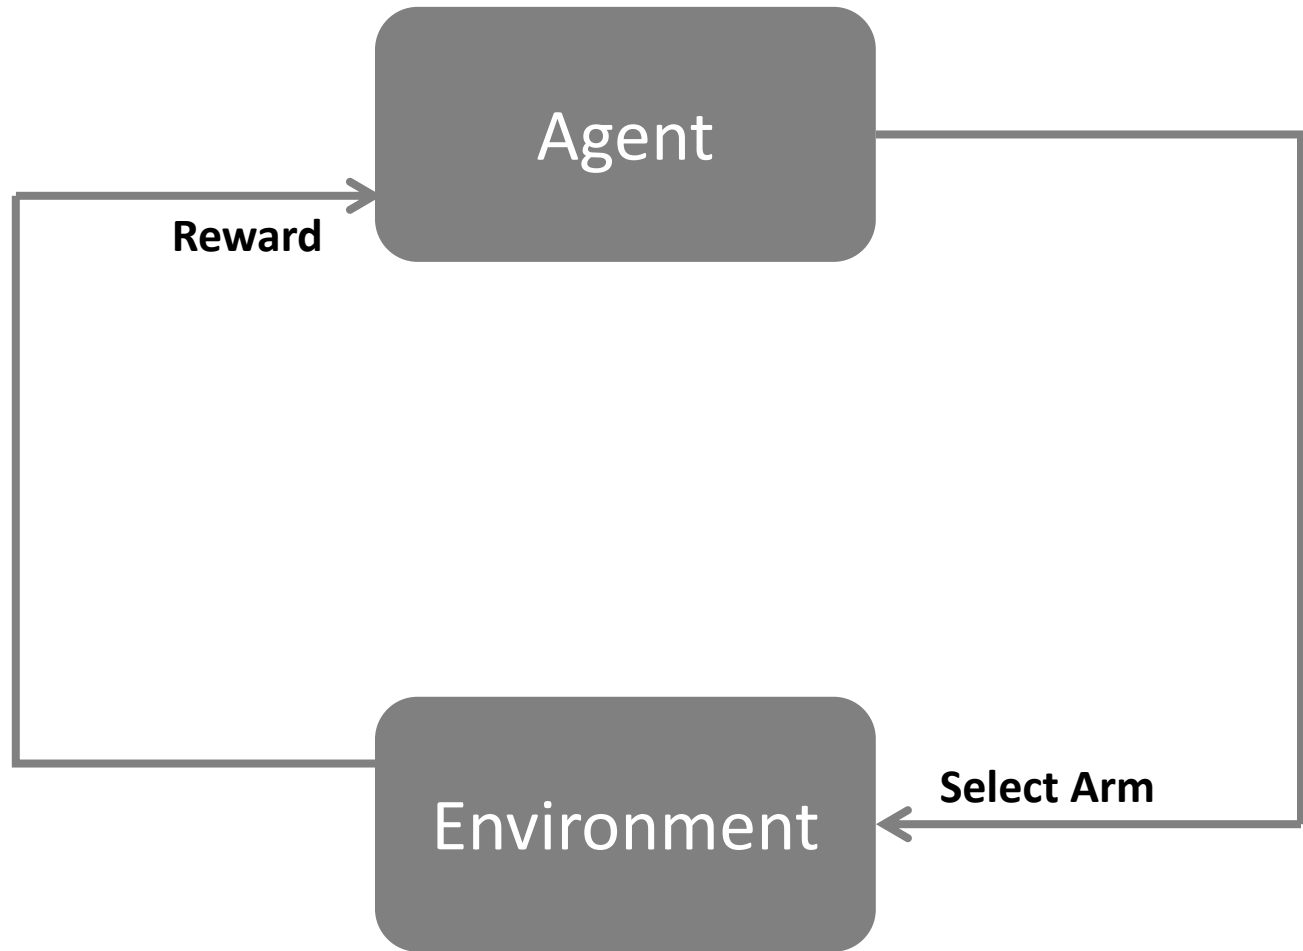
Bandit/Arm 3

78 %

**current**  
success (win)  
rate

Which bandit shall we play next?

# K-Armed Bandit



# Exploration vs. Exploitation

The crucial tradeoff the gambler faces at each trial is between "**exploitation**" of the machine that has the highest expected payoff and "**exploration**" to get more information about the expected payoffs of the other machines.

# $\epsilon$ -greedy Algorithm

generate random number  $p \in [0, 1]$

if ( $p < \epsilon$ )           // explore

    select random arm


else                   // exploit

    select current best arm

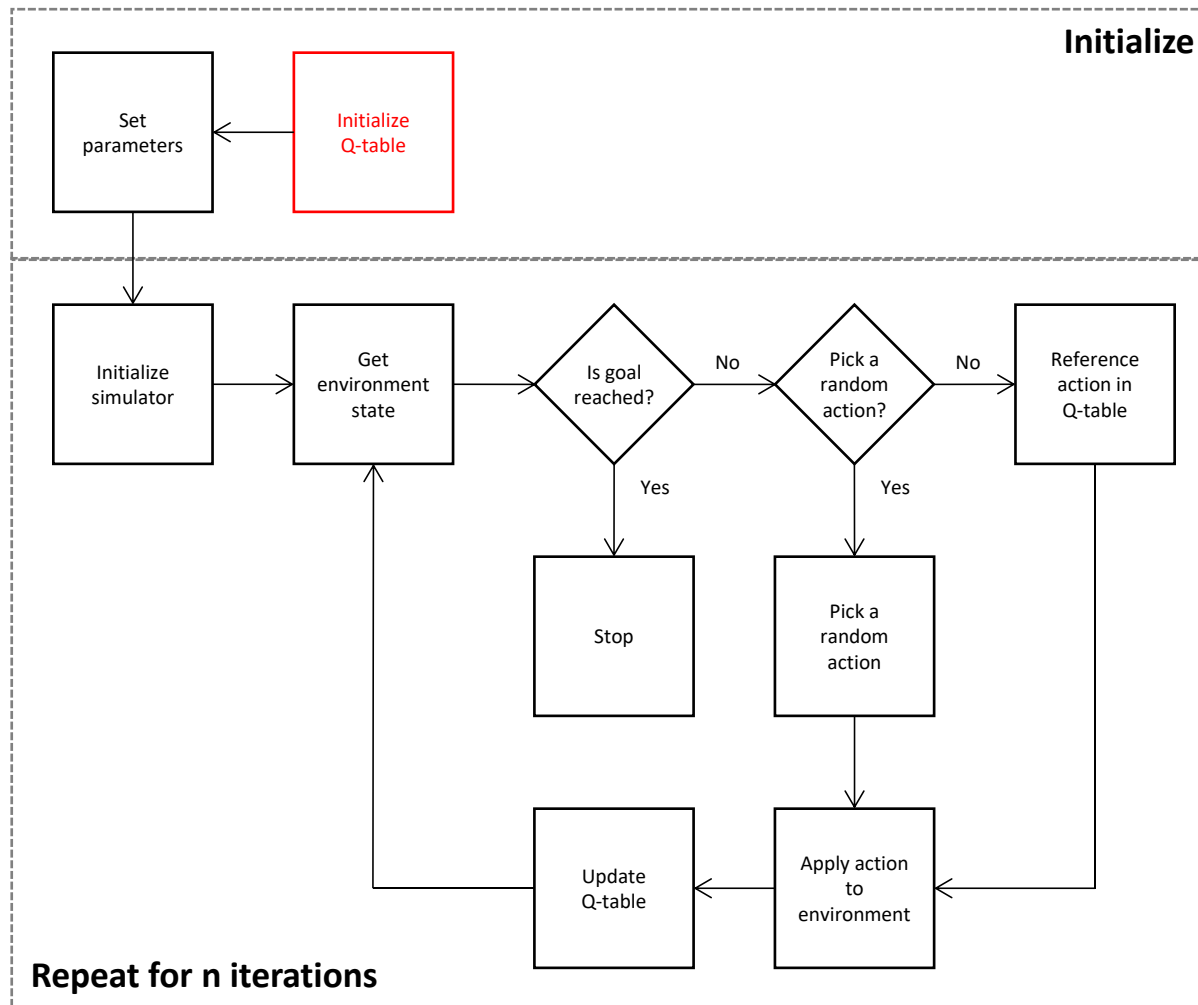
end



# Q-Learning Algorithm



# Q-Learning Algorithm



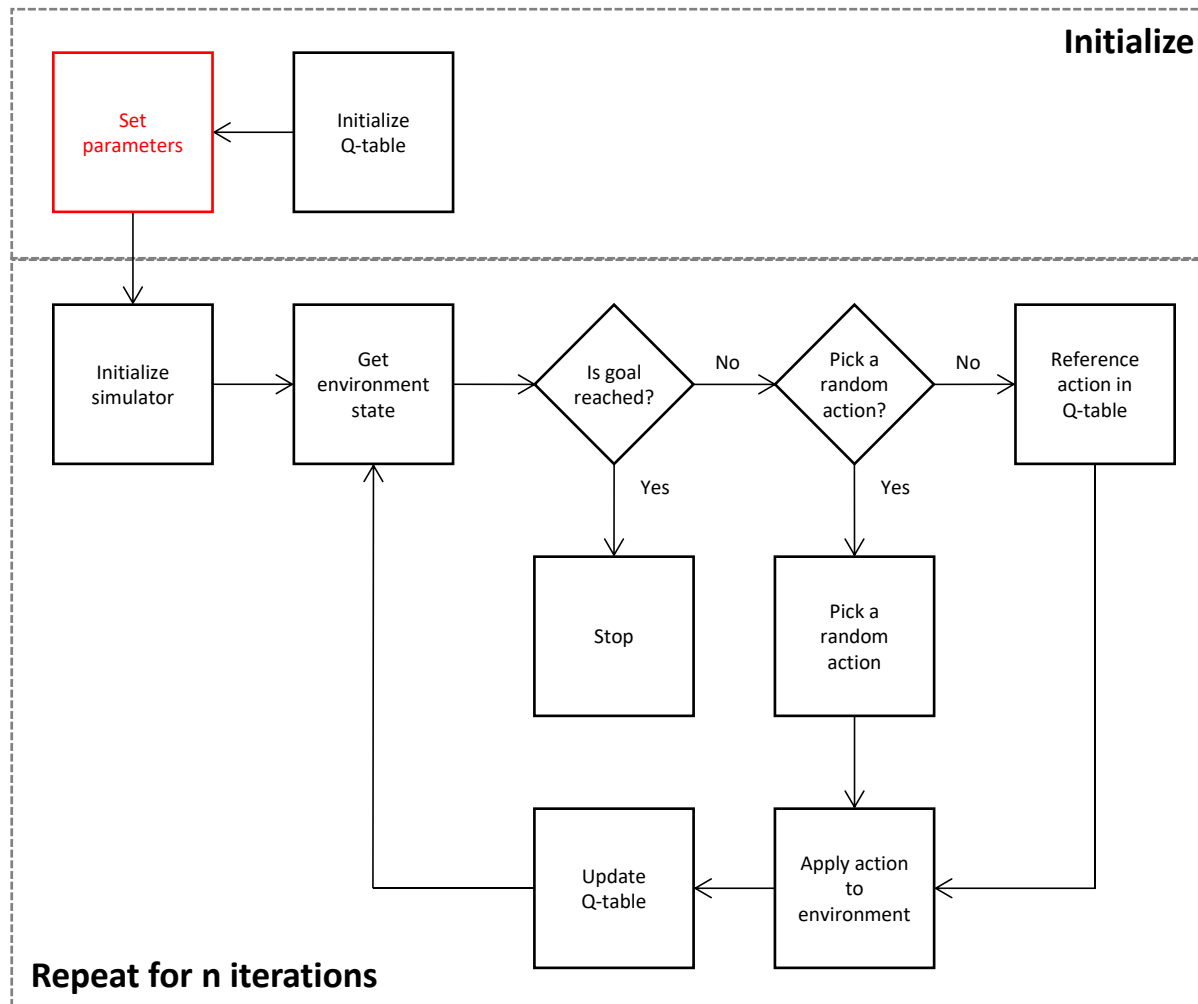
## Initialize Q-table:

Set up and initialize (all values set to 0) a table where:

- rows represent **possible states**
- columns represent **actions**

Note that additional states can be added to the table when encountered.

# Q-Learning Algorithm



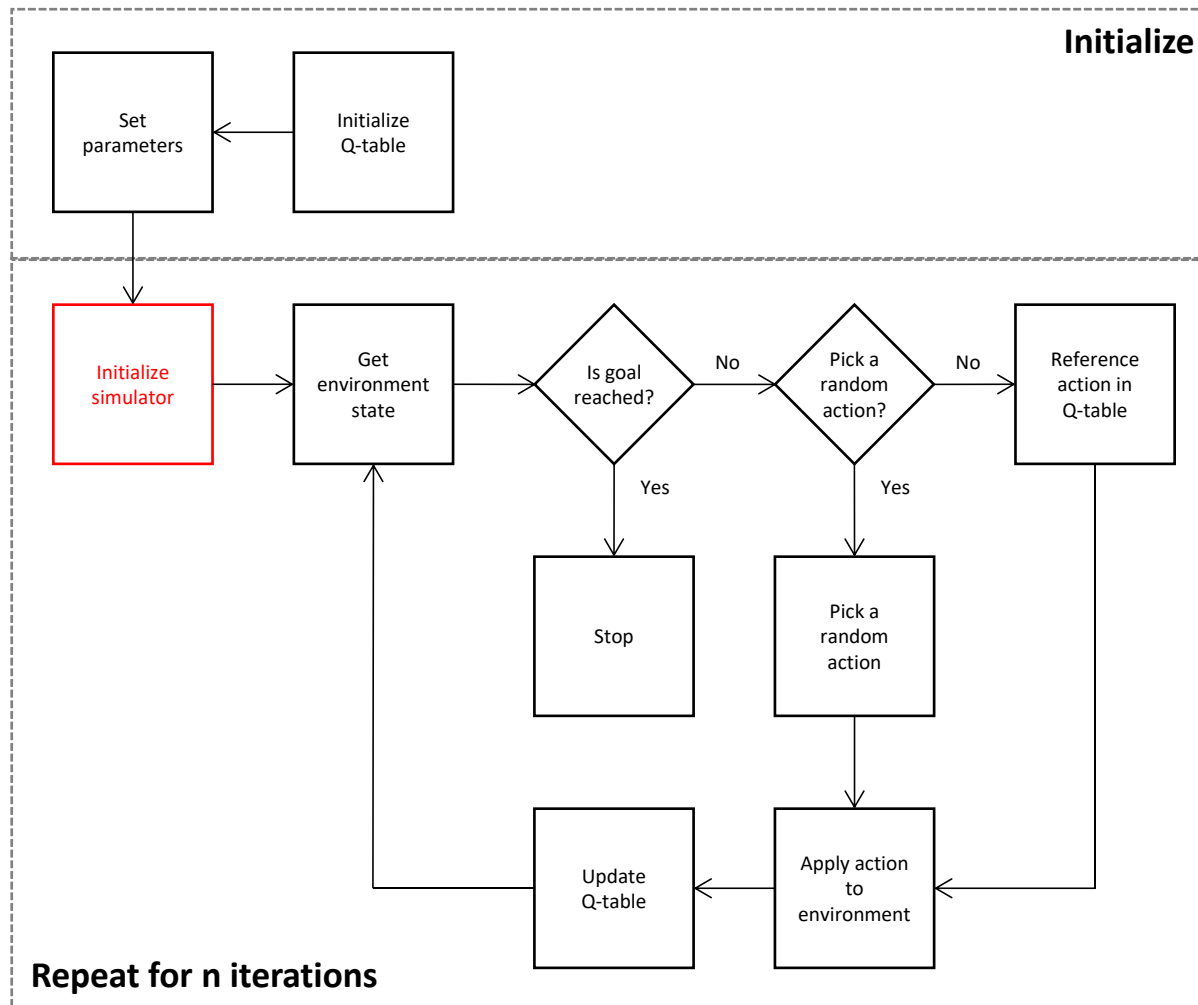
## Set parameters:

Set and initialize **hyperparameters** for the Q-learning process.

## Hyperparameters include:

- **chance of choosing a random action:** a threshold for **choosing a random action over an action from the Q-table**
- **learning rate:** a parameter that describes **how quickly the algorithm should learn from rewards** in different states
  - high: faster learning with erratic Q-table changes
  - low: gradual learning with possibly more iterations
- **discount factor:** a parameter that **describes how valuable are future rewards**. It tells the algorithm whether it should seek “immediate gratification” (small) or “long-term reward” (large)

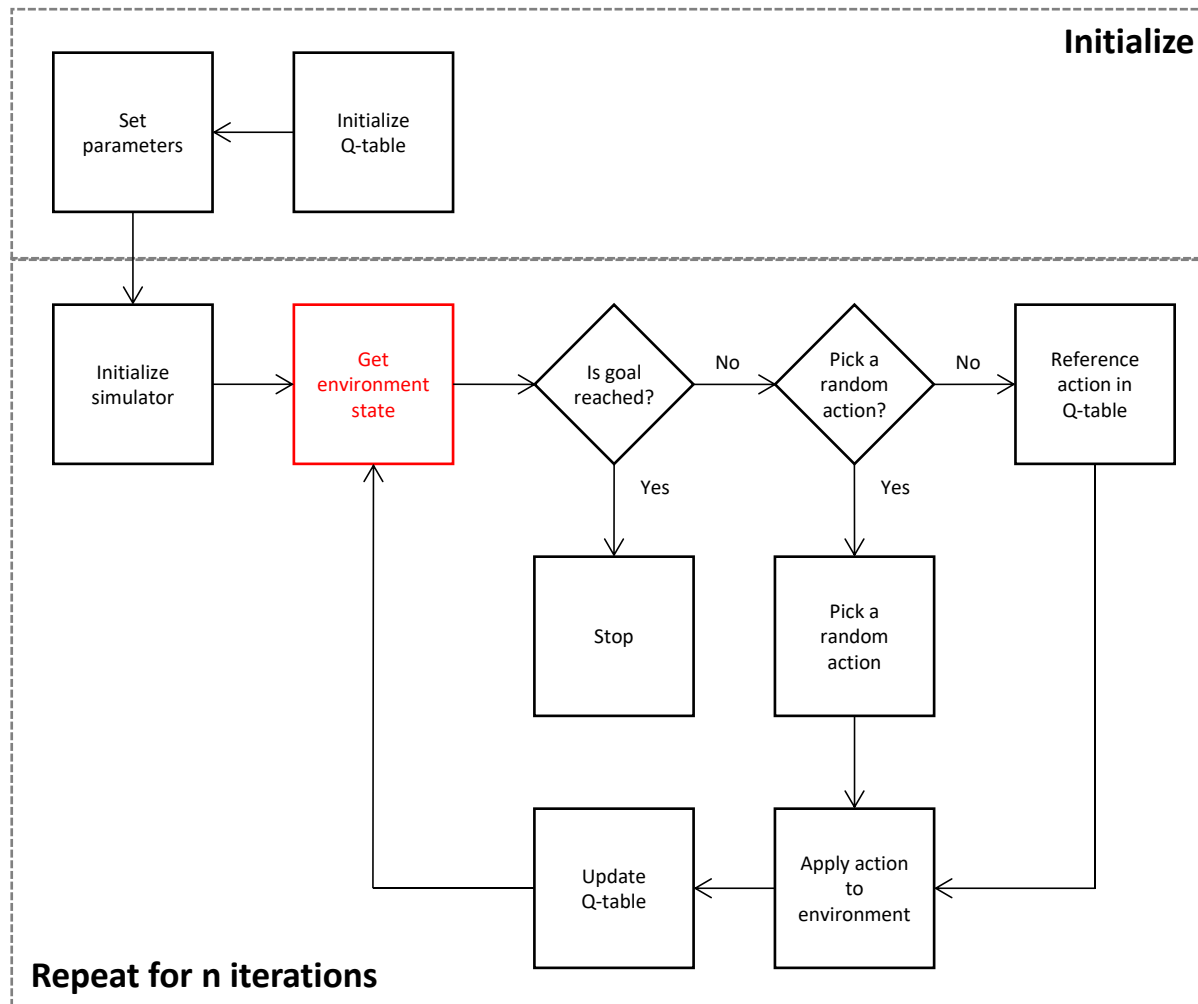
# Q-Learning Algorithm



## Initialize simulator:

Reset the simulated environment to its initial state and place the agent in a neutral state.


# Q-Learning Algorithm



## Get environment state:

Report the current state of the environment. Typically a vector of values representing all relevant variables.

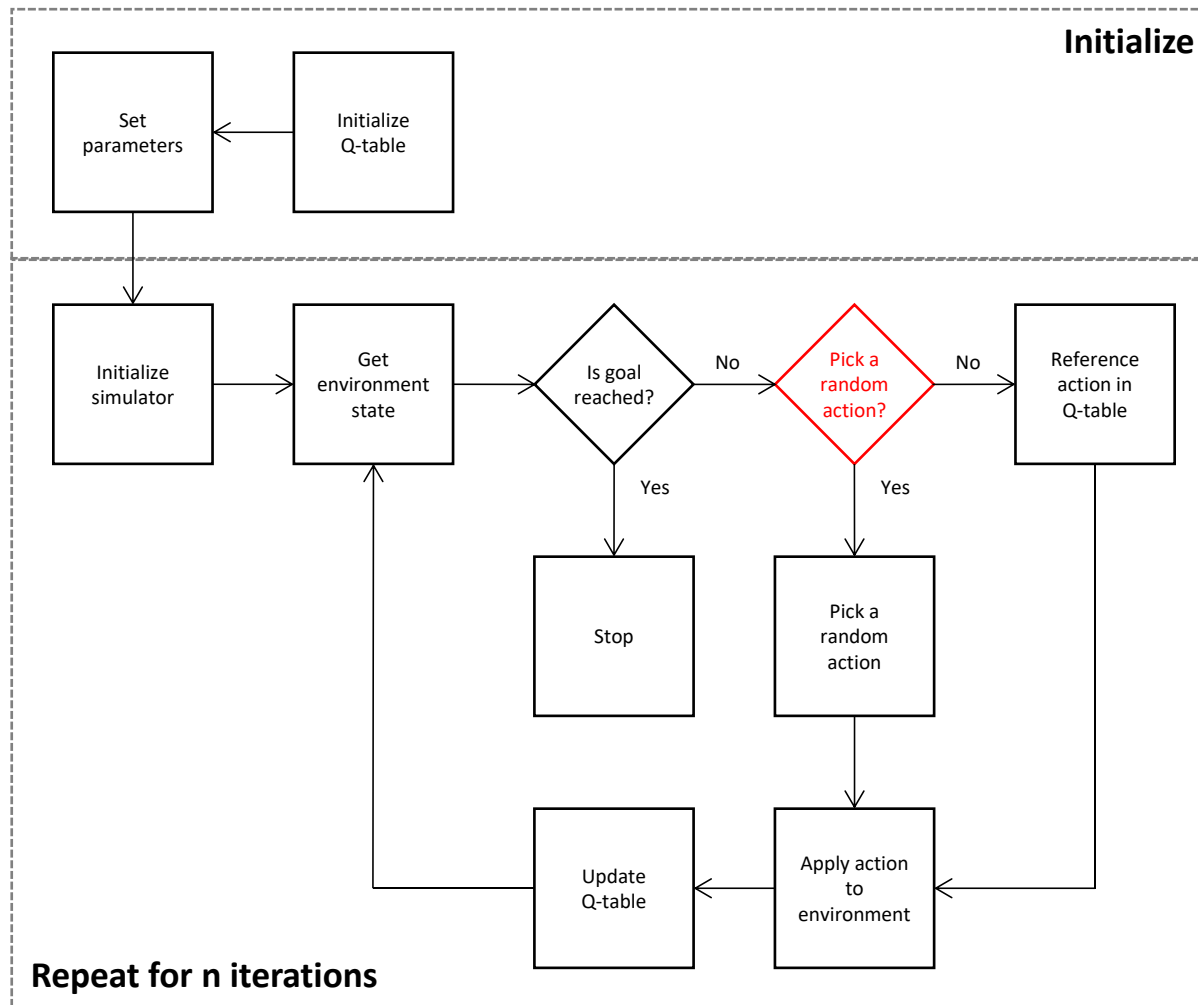
# Q-Learning Algorithm



## Is goal reached?:

Verify if the goal of the simulation has been achieved. It could be decided with the agent arriving in expected final state or by some simulation parameter.

# Q-Learning Algorithm



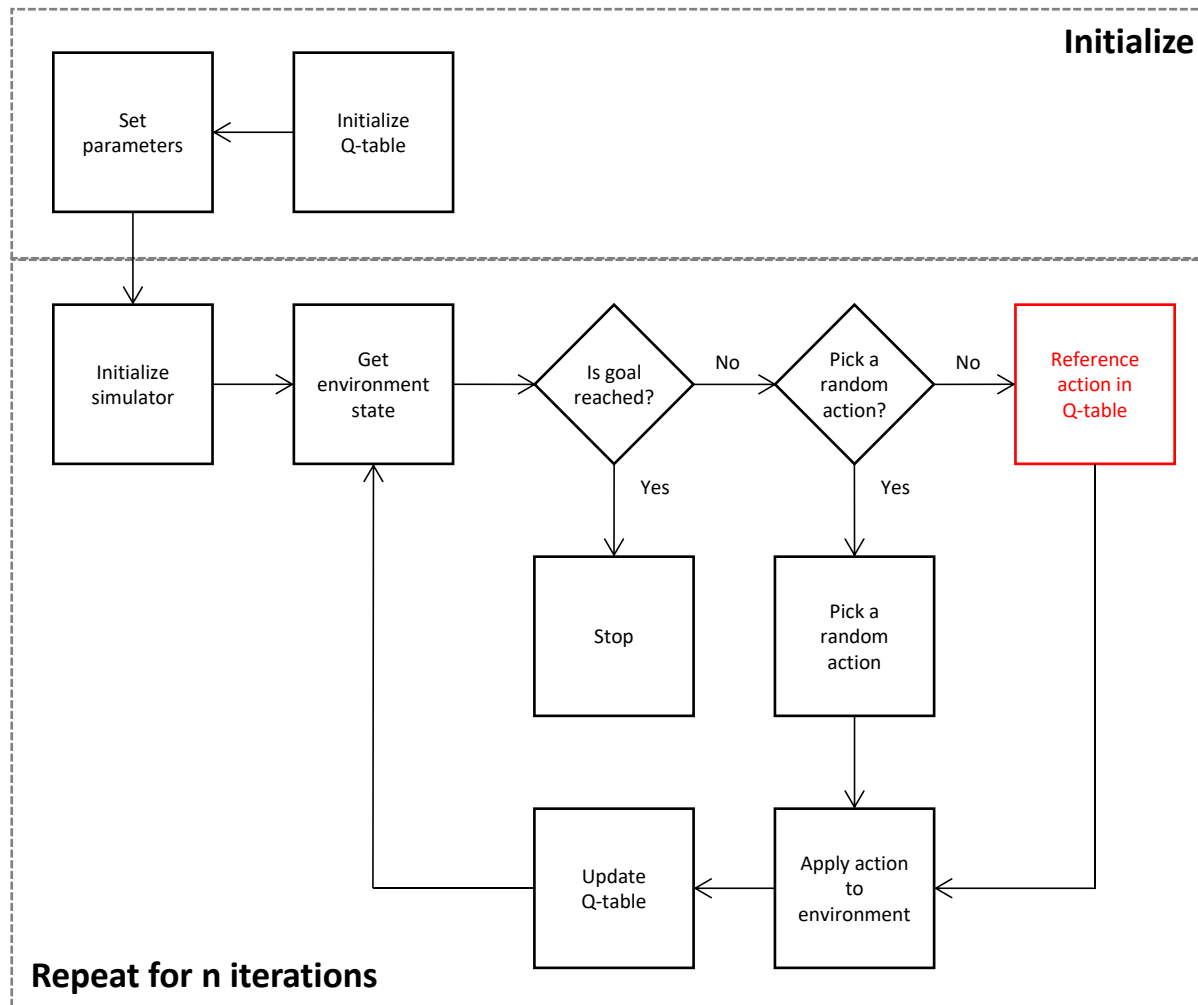
## Pick a random action?:

Decide whether next action should be picked at random or not (it will be selected based on Q-table data then).

Use the **chance of choosing a random action hyperparameter** to decide.



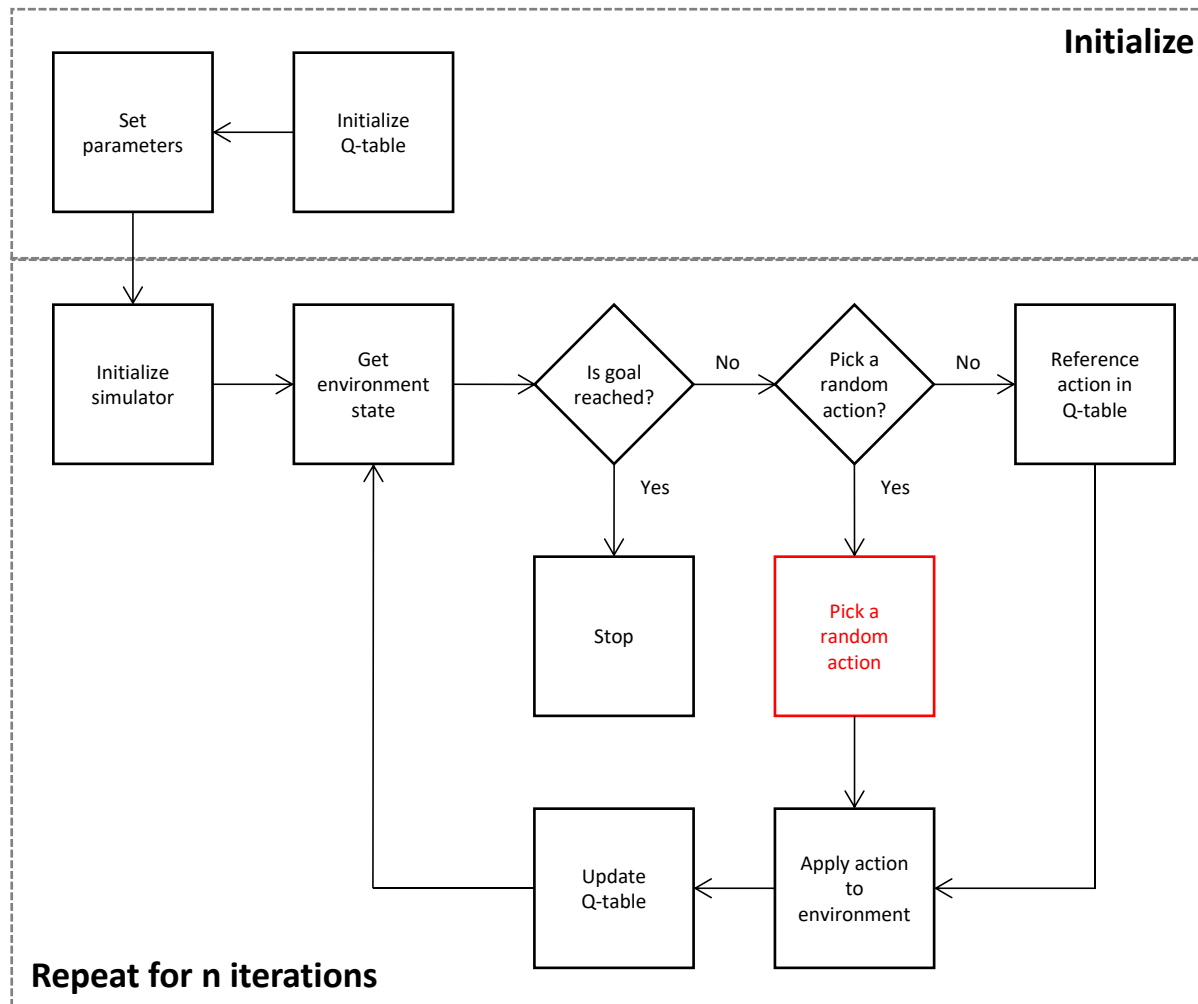
# Q-Learning Algorithm



**Reference action in Q-table:**

Next action decision will be based on data from the Q-table **given the current state of the environment.**

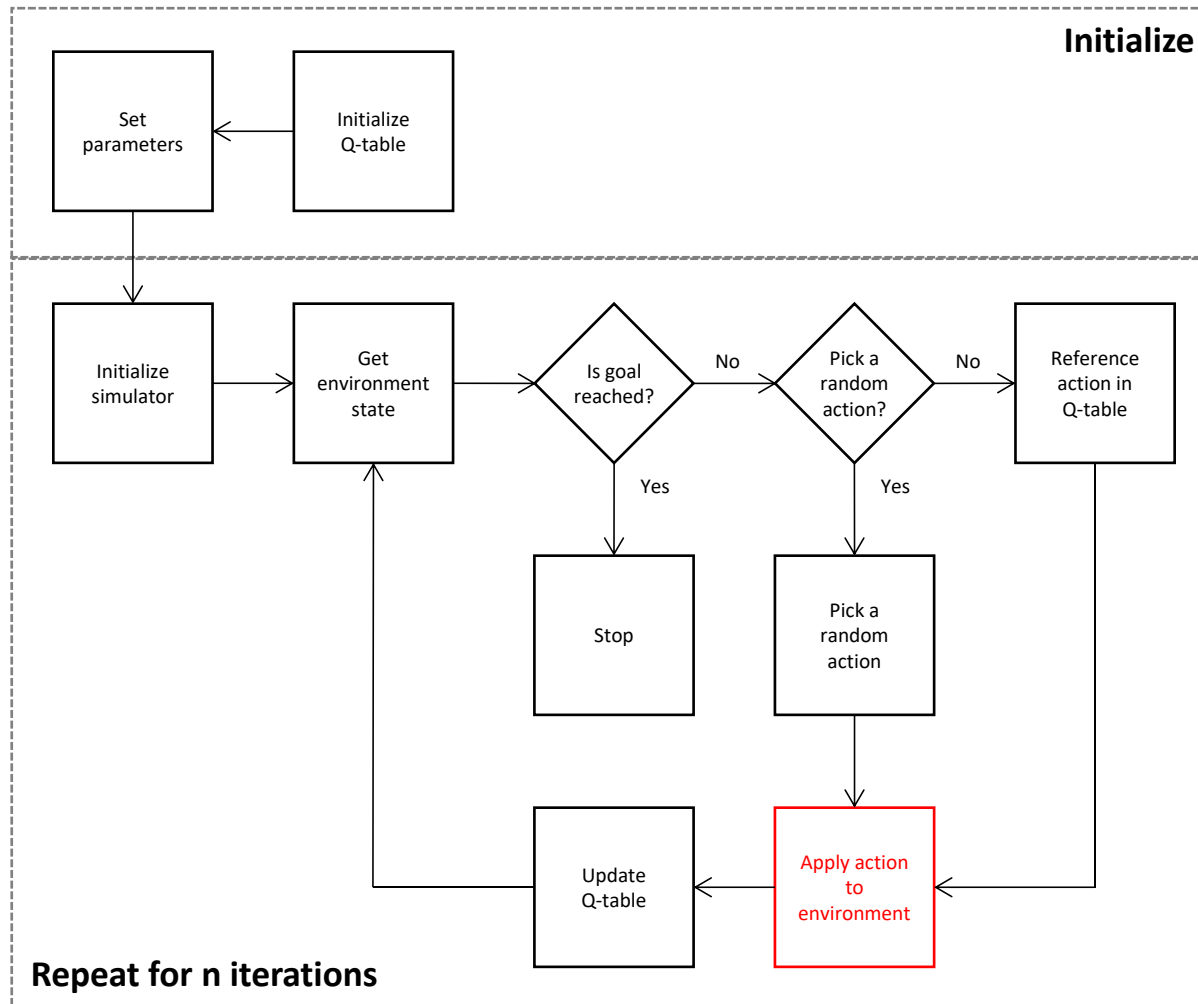
# Q-Learning Algorithm



## Pick a random action:

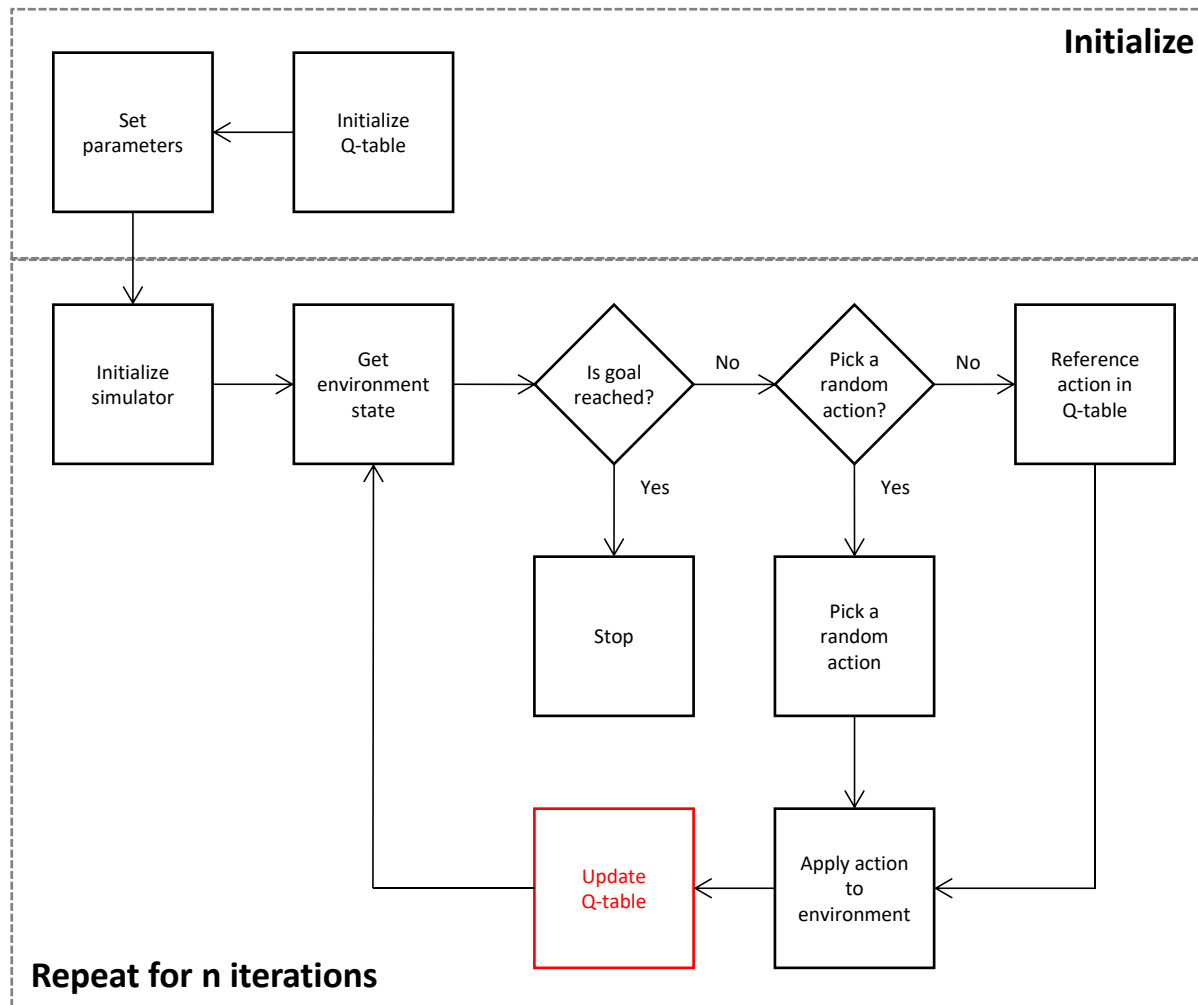
Pick any of the available actions at random. Helpful with exploration of the environment.

# Q-Learning Algorithm



**Apply action to environment:**  
Apply the action to the environment to change it. Each action will have its own reward.

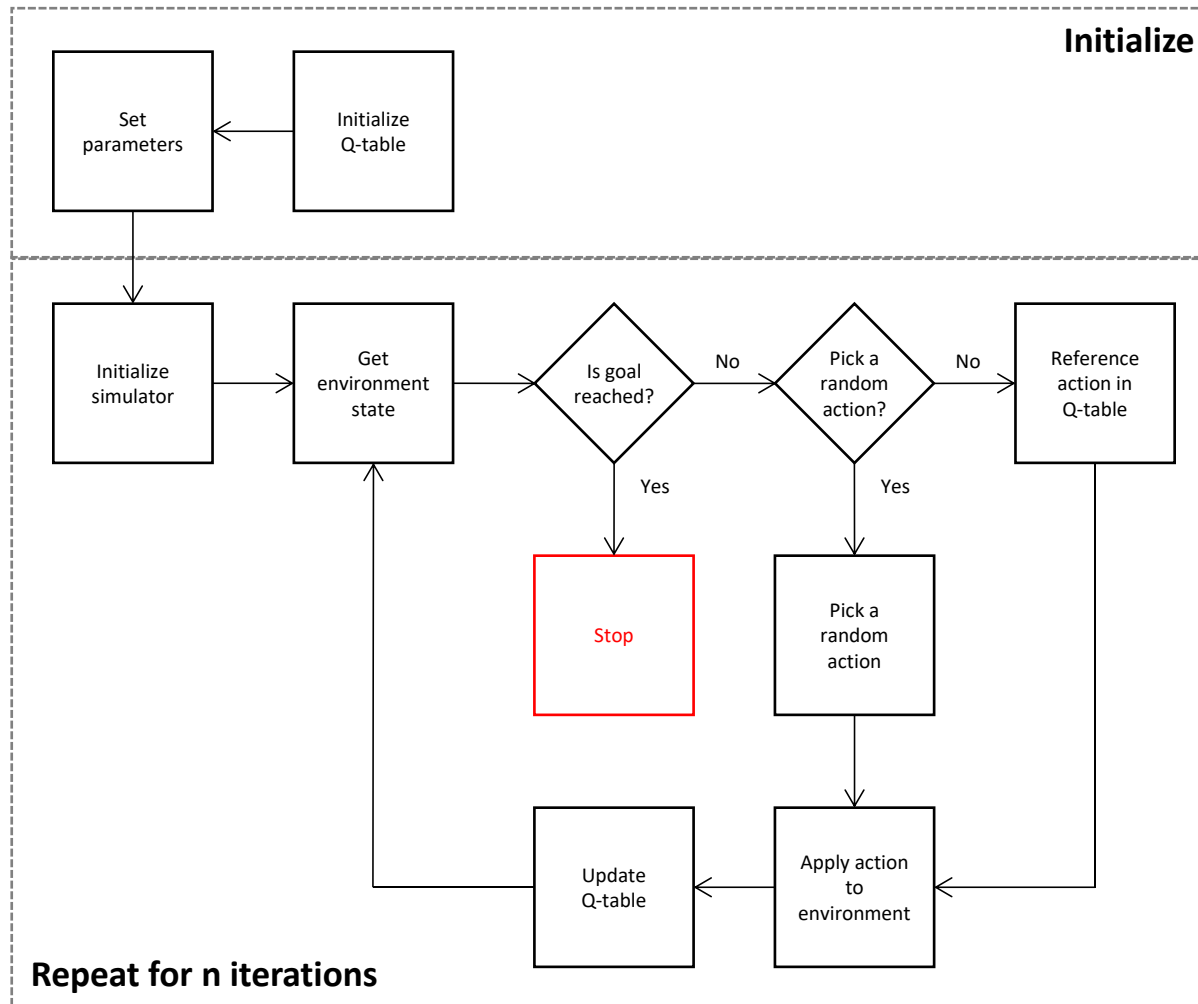
# Q-Learning Algorithm



## Update Q-table:


Update the Q-table given the reward resulting from recently applied action (feedback from the environment).

# Q-Learning Algorithm



**Stop:**  
Stop the learning process

# Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0
	...	...	...	...	...
	n	0	0	0	0

## Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:


Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

← Learning rate
Discount

Current value
Maximum value of all actions on next state

# Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0
	...	...	...	...	...
	n	0	0	0	0

## Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action:

Reward:

Q-table value:


$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

← Learning rate
Discount

Current value
Maximum value of all actions on next state →



# Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0
	...	...	...	...	...
	n	0	0	0	0

### Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

## Action:

## Reward:


### Q-table value:

$$Q(\text{state}, \text{action}) = (1 - \text{alpha}) * Q(\text{state}, \text{action}) + \text{alpha} * (\text{reward} + \text{gamma} * Q(\text{next state}, \text{all actions}))$$

Learning rate      Discount

Current value      Maximum value of all actions on next state

# Q-Learning Algorithm



Action: →

Reward: -100

Q-table value:

$$Q(1, \text{east}) = (1 - 0.1) * 0 + 0.1 * (-100 + 0.6 * \max \text{ of } Q(2, \text{all actions}))$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	0	0
	2	0	0	0	0
	...	...	...	...	...
	n	0	0	0	0

## Rewards:















































Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

# Q-Learning Algorithm

 → 								
								
								
								
								
								
								
								
								
								

Action: →

Reward:   -100

Q-table value:

$$Q(1, \text{east}) = (1 - 0.1) * 0 + 0.1 * (-100 + 0.6 * 0) = -10$$

Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	0	0	0
	...	...	...	...	...
	n	0	0	0	0

**Rewards:**


Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

# Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	0	0	0
	...	...	...	...	...
	n	0	0	0	0

## Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500


Action: →

Reward: -1000

Q-table value:

$$Q(2, \text{south}) = (1 - 0.1) * 0 + 0.1 * (-1000 + 0.6 * \max \text{ of } Q(3, \text{all actions}))$$

# Q-Learning Algorithm



Q-table		Actions			
		↑	↓	→	←
States	1	0	0	-10	0
	2	0	-100	0	0
	...	...	...	...	...
	n	0	0	0	0

## Rewards:

Move into car: -100

Move into pedestrian: -1000

Move into empty space: 100

Move into goal: 500

Action: →

Reward:   -1000

Q-table value:

$$Q(2, \text{south}) = (1 - 0.1) * 0 + 0.1 * (-1000 + 0.6 * 0) = -100$$

# Deep Reinforcement Learning

