

# Maximum entropy models and structured prediction

CS-585

Natural Language Processing

Derrick Higgins

# HMMs, MEMMs and CRFs

---

- Hidden Markov Models
  - Generative sequence labeling models
- Maximum Entropy Markov Models
  - Like HMMs, but discriminative
  - Logistic regression
- Conditional Random Fields
  - Like MEMMs, but in structured prediction paradigm

---

# MAXIMUM ENTROPY MARKOV MODELS

# Maximum Entropy

---

- In NLP, logistic regression models are often called *maximum entropy* (or *maxent*) models
- Idea
  - when selecting a probability distribution to model observed data, we want the distribution to model the statistics of the data, but otherwise reserve judgement
  - Model should be as uncertain as possible, while still capturing the data
  - Uncertainty=entropy, so select the model with maximal entropy subject to data-fitting constraints



# Maximum Entropy

- In a supervised learning context, the data fitting constraints have to do with the frequencies of feature functions  $f_k(X, Y)$  in the training data:

$$\sum_n f_k(X_n, Y_n) P(X_n, Y_n) = c_k$$

- It can be shown that the distribution  $P(Y_n|X_n)$  with maximum entropy subject to these constraints has the form

$$P(Y_n|X_n) = \frac{e^{\sum_k \lambda_k f_k(X_n, Y_n)}}{Z}$$

- Where  $Z$  is a normalizing constant
- Just logistic regression....

# Maximum Entropy Markov Models (MEMMs)

- Cousins of HMMs

- Still based on the Markov assumption:

$$P(T|W) = \prod_i P(T_i | T_{i-k..i-1}, W)$$

- Because they are discriminative, MEMMs cannot be trained in an unsupervised way like HMMs
  - HMMs are generative (model  $P(W, T)$ )
  - MEMMs are discriminative (model  $P(T|W)$ )
- Based on maximum entropy (logistic regression) models to predict tag for each word given local context
  - Can be trained using gradient descent or model-specific algorithms (GIS, IIS)

# Maximum Entropy Markov Models (MEMMs)

Estimate conditional probability of tags given text

$$P_m(T|W) = \prod_i P(T_i | T_{i-k..i-1}, W)$$

(Markov assumption)

$$= \prod_i \frac{e^{\sum_k \lambda_k f_k(W, T_{i-k..i})}}{\sum_{t \in \mathcal{T}} e^{\sum_k \lambda_k f_k(W, T_{i-k..i-1}, t)}}$$

(maxent)

- Can be trained using gradient descent
  - Either using forward-backward algorithm and likelihood (regularized maximum likelihood)
- Similar dynamic programming tricks to HMMs for efficiently computing sum across all labelings

**Locally normalized**

# MEMMs: POS feature functions

- Remember Transformation-Based Learning

The preceding (following) word is tagged **z**.  
The word two before (after) is tagged **z**.  
One of the two preceding (following) words is tagged **z**.  
One of the three preceding (following) words is tagged **z**.  
The preceding word is tagged **z** and the following word is tagged **w**.  
The preceding (following) word is tagged **z** and the word two before (after) is tagged **w**.

Change tags		Condition	Example
#	From To		
1	NN VB	Previous tag is TO	to/TO race/NN → VB
2	VBP VB	One of the previous 3 tags is MD	might/MD vanish/VBP → VB
3	NN VB	One of the previous 2 tags is MD	might/MD not reply/NN → VB
4	VB NN	One of the previous 2 tags is DT	
5	VBD VBN	One of the previous 3 tags is VBZ	



# MEMMMs: POS feature functions

- Maxent feature functions

Condition	Features
$w_i$ is not rare	$w_i = X$ & $t_i = T$
$w_i$ is rare	$X$ is prefix of $w_i$ , $ X  \leq 4$ & $t_i = T$
	$X$ is suffix of $w_i$ , $ X  \leq 4$ & $t_i = T$
	$w_i$ contains number & $t_i = T$
	$w_i$ contains uppercase character & $t_i = T$
	$w_i$ contains hyphen & $t_i = T$
$\forall w_i$	$t_{i-1} = X$ & $t_i = T$
	$t_{i-2}t_{i-1} = XY$ & $t_i = T$
	$w_{i-1} = X$ & $t_i = T$
	$w_{i-2} = X$ & $t_i = T$
	$w_{i+1} = X$ & $t_i = T$
	$w_{i+2} = X$ & $t_i = T$

Table 1: Features on the current history  $h_i$

# MEMMs: POS feature functions

Word	the	stories	about	well-heeled	communities	and	developers
Tag	DT	NNS	IN	JJ	NNS	CC	NNS
Position	1	2	3	4	5	6	7

$w_i = \text{about}$	$\wedge t_i = \text{IN}$
$w_{i-1} = \text{stories}$	$\wedge t_i = \text{IN}$
$w_{i-2} = \text{the}$	$\wedge t_i = \text{IN}$
$w_{i+1} = \text{well-heeled}$	$\wedge t_i = \text{IN}$
$w_{i+2} = \text{communities}$	$\wedge t_i = \text{IN}$
$t_{i-1} = \text{NNS}$	$\wedge t_i = \text{IN}$
$t_{i-2}t_{i-1} = \text{DT NNS}$	$\wedge t_i = \text{IN}$

# MEMMs: POS feature functions

Word	the	stories	about	well-heeled	communities	and	developers
Tag	DT	NNS	IN	JJ	NNS	CC	NNS
Position	1	2	3	4	5	6	7

$prefix(w_i) = w$	$\wedge t_i = JJ$		
$prefix(w_i) = we$	$\wedge t_i = JJ$		
$prefix(w_i) = wel$	$\wedge t_i = JJ$		
$prefix(w_i) = well$	$\wedge t_i = JJ$		
$suffix(w_i) = d$	$\wedge t_i = JJ$		
$suffix(w_i) = ed$	$\wedge t_i = JJ$		
$suffix(w_i) = led$	$\wedge t_i = JJ$		
$suffix(w_i) = eled$	$\wedge t_i = JJ$		
$w_i$ contains hyphen	$\wedge t_i = JJ$		
		$w_{i-1} = \text{about}$	$\wedge t_i = JJ$
		$w_{i-2} = \text{stories}$	$\wedge t_i = JJ$
		$w_{i+1} = \text{communities}$	$\wedge t_i = JJ$
		$w_{i+2} = \text{and}$	$\wedge t_i = JJ$
		$t_{i-1} = \text{IN}$	$\wedge t_i = JJ$
		$t_{i-2}t_{i-1} = \text{NNS IN}$	$\wedge t_i = JJ$

# Viterbi Tagging for MEMMs

Most probable tag sequence given text:

$$T^* = \operatorname{argmax}_T P_m(T|W)$$

$$= \operatorname{argmax}_T \prod_i P(T_i | T_{i-k..i-1}, W)$$


(Markov assumption)

$$= \operatorname{argmax}_T \prod_i \frac{e^{\sum_k \lambda_k f_k(W, T_{i-k..i})}}{\sum_{t \in \mathcal{T}} e^{\sum_k \lambda_k f_k(W, T_{i-k..i-1}, t)}}$$

(maxent)

$$= \operatorname{argmax}_T \sum_i \left( \sum_k \lambda_k f_k(W, T_{i-k..i}) - \log \sum_{t \in \mathcal{T}} e^{\sum_k \lambda_k f_k(W, T_{i-k..i-1}, t)} \right)$$


# Viterbi algorithm



$$-\log P(t_i | t_{i-1}, w_i, w_{i-1})$$

t <sub>i-1</sub>	t <sup>0</sup>		t <sup>1</sup>									t <sup>2</sup>								
w <sub>i</sub>	w <sup>1</sup>	w <sup>2</sup>	w <sup>1</sup>			w <sup>2</sup>			w <sup>3</sup>			w <sup>1</sup>			w <sup>2</sup>			w <sup>3</sup>		
w <sub>i-1</sub>			w <sup>1</sup>	w <sup>2</sup>	w <sup>3</sup>	w <sup>1</sup>	w <sup>2</sup>	w <sup>3</sup>	w <sup>1</sup>	w <sup>2</sup>	w <sup>3</sup>	w <sup>1</sup>	w <sup>2</sup>	w <sup>3</sup>	w <sup>1</sup>	w <sup>2</sup>	w <sup>3</sup>	w <sup>1</sup>	w <sup>2</sup>	w <sup>3</sup>
t <sub>i</sub> =t <sup>1</sup>	1.0	1.1	2	4	2	2	3	2	4	5	2	2	2	3	2	4	5	2	4	
t <sub>i</sub> =t <sup>2</sup>	1.1	2.1	3	5	3	3	2	4	2	4	3	3	3	2	4	2	4	3	4	5


# Viterbi algorithm



$$-\log P(t_i | t_{i-1}, w_i, w_{i-1})$$

$t_{i-1}$	$t^0$		$t^1$									$t^2$								
$w_i$	$w^1$	$w^2$	$w^1$			$w^2$			$w^3$			$w^1$			$w^2$			$w^3$		
$w_{i-1}$			$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$
$t_i = t^1$	1.0	1.1	2	4	2	2	2	3	2	4	5	2	2	2	3	2	4	5	2	4
$t_i = t^2$	1.1	2.1	3	5	3	3	2	4	2	4	3	3	3	2	4	2	4	3	4	5

# Viterbi algorithm



$$-\log P(t_i | t_{i-1}, w_i, w_{i-1})$$

$t_{i-1}$	$t^0$		$t^1$									$t^2$								
$w_i$	$w^1$	$w^2$	$w^1$			$w^2$			$w^3$			$w^1$			$w^2$			$w^3$		
$w_{i-1}$			$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$
$t_i = t^1$	1.0	1.1	2	4	2	2	2	3	2	4	5	2	2	2	3	2	4	5	2	4
$t_i = t^2$	1.1	2.1	3	5	3	3	2	4	2	4	3	3	3	2	4	2	4	3	4	5

# Viterbi Algorithm

$D(0, \text{START}) = 0$

**for each** tag  $t \neq \text{START}$  **do:**

$D(0, t) = -\infty$

**for**  $i \leftarrow 1$  **to**  $N$  **do:**

**for each** tag  $t^j$  **do:**

$$D(i, t^j) \leftarrow \max_k (D(i-1, t^k) + \log P(t_i | t_{i-1}=t^k, W))$$

$\log P(T|W) = \max_j D(N, t^j)$




# MEMM Limitations

---

- Locally normalized
  - MEMM assumes that the probability distribution over a tagging  $T$  can be factored into the product of conditional probabilities with limited history for each tag location in a sentence:
  - But this limits the flexibility of the model. Only paths from the same prior state can “compete” against one another for probability mass

# MEMM Limitations: Label Bias

	PERSON count	LOCATION count
Harvey Ford	9	1
Harvey Park	1	9
Myrtle Ford	9	1
Myrtle Park	1	9



# MEMM Limitations: Label Bias

	PERSON count	LOCATION count
Harvey Ford	9	1
Harvey Park	1	9
Myrtle Ford	9	1
Myrtle Park	1	9

Conditional probabilities:

$$P(t_i = \text{B-PER} | t_{i-1} = \text{Start}, w_i = \text{Harvey}) = 0.5$$

$$P(t_i = \text{B-LOC} | t_{i-1} = \text{Start}, w_i = \text{Harvey}) = 0.5$$

$$P(t_i = \text{B-PER} | t_{i-1} = \text{Start}, w_i = \text{Myrtle}) = 0.5$$

$$P(t_i = \text{B-LOC} | t_{i-1} = \text{Start}, w_i = \text{Myrtle}) = 0.5$$

$$P(t_i = \text{I-PER} | t_{i-1} = \text{B-PER}, w_i = \text{Ford}) = 1.0$$

$$P(t_i = \text{I-LOC} | t_{i-1} = \text{B-LOC}, w_i = \text{Ford}) = 1.0$$

$$P(t_i = \text{I-PER} | t_{i-1} = \text{B-PER}, w_i = \text{Park}) = 1.0$$

$$P(t_i = \text{I-LOC} | t_{i-1} = \text{B-LOC}, w_i = \text{Park}) = 1.0$$

# MEMM Limitations: Label Bias

0.5      1.0  
B-LOC   E-LOC

Harvey   Park

0.5      1.0  
B-PER   E-PER

Harvey   Park



Conditional probabilities:

$$P(t_i = \text{B-PER} | t_{i-1} = \text{Start}, w_i = \text{Harvey}) = 0.5$$

$$P(t_i = \text{B-LOC} | t_{i-1} = \text{Start}, w_i = \text{Harvey}) = 0.5$$

$$P(t_i = \text{B-PER} | t_{i-1} = \text{Start}, w_i = \text{Myrtle}) = 0.5$$

$$P(t_i = \text{B-LOC} | t_{i-1} = \text{Start}, w_i = \text{Myrtle}) = 0.5$$

$$P(t_i = \text{I-PER} | t_{i-1} = \text{B-PER}, w_i = \text{Ford}) = 1.0$$

$$P(t_i = \text{I-LOC} | t_{i-1} = \text{B-LOC}, w_i = \text{Ford}) = 1.0$$

$$P(t_i = \text{I-PER} | t_{i-1} = \text{B-PER}, w_i = \text{Park}) = 1.0$$

$$P(t_i = \text{I-LOC} | t_{i-1} = \text{B-LOC}, w_i = \text{Park}) = 1.0$$

Really? But this was 9 times more common in the training data!

# MEMM Limitations: Label Bias

---

- *Label bias* problem: Low-entropy states (from which following states are highly predictable) play an unreasonably strong role in determining label sequence
  - In an extreme case, causing the words to be irrelevant in determining the labels
- Due to local normalization of probabilities for each tag, instead of global normalization for entire tag sequence

---

# CONDITIONAL RANDOM FIELDS

# Structured prediction

---

- CRFs fall into a predictive modeling framework called *structured prediction*
  - When we have some complex structured object over which we want to make predictions...
    - pixels within an image, tag sequences or syntactic trees for a text
  - ... we try to estimate a probability distribution over the whole output space, rather than distributions over subparts

# Structured prediction

Structured prediction for sequence labeling: predict optimal tagging  $T^*$  using a scoring function over the output space:

$$\hat{T} = \operatorname{argmax}_T \Psi(T, W)$$

In a probabilistic framework:

$$\hat{T} = \operatorname{argmax}_T P(T|W)$$

Specifically using logistic regression:

$$\hat{T} = \operatorname{argmax}_T \frac{\prod_i e^{\lambda_i f_i(T, W)}}{\sum_{T'} \prod_i e^{\lambda_i f_i(T', W)}}$$



# Feature functions for CRFs

---

- Essentially the same as for MEMMs: can incorporate left and right context for observations (words)
- For tags, the probabilistic framework theoretically could encompass features built on arbitrary tag dependencies
  - E.g., first and last tag in sentence
  - But we need dynamic programming to make inference tractable. Therefore, in practice, tag dependencies are limited to adjacent tags
- “Linear chain CRF”

# The partition function

- Remember our expression for the probability of a tag sequence under a CRF:

$$P(T|W) = \frac{\prod_i e^{\lambda_i f_i(T,W)}}{\sum_{T'} \prod_i e^{\lambda_i f_i(T',W)}}$$

The denominator is called the partition function. It involves a sum over all possible tag sequences for the sentence and is expensive to compute

But note that we don't need to compute it in order to predict the best tagging; the numerator is sufficient

# Inference in CRFs: the Viterbi algorithm

---

$$\begin{aligned}\operatorname{argmax}_T P(T|W) &= \operatorname{argmax}_T \frac{\prod_i e^{\lambda_i f_i(T,W)}}{\sum_{T'} \prod_i e^{\lambda_i f_i(T',W)}} \\ &= \operatorname{argmax}_T \prod_i e^{\lambda_i f_i(T,W)} \\ &= \operatorname{argmax}_T \sum_i \lambda_i f_i(T,W)\end{aligned}$$

# Inference in CRFs: the Viterbi algorithm

---

$$\begin{aligned}\operatorname{argmax}_T P(T|W) &= \operatorname{argmax}_T \sum_i \lambda_i f_i(T, W) \\ &= \operatorname{argmax}_T \sum_{j=1}^N \sum_k \lambda_k f_k(T_{j-1}, T_j, W)\end{aligned}$$


Because feature functions only express dependencies between adjacent tag pairs, we can group them into sets based on the final tag in the dependency chain

# Inference in CRFs: the Viterbi algorithm

Word	the	stories	about	well-heeled	communities	and	developers
Tag	DT	NNS	IN	JJ	NNS	CC	NNS
Position	1	2	3	4	5	6	7

$$f_{1..K}(T_1, T_2, T_3, T_4, T_5, T_6, T_7, W)$$


# Viterbi algorithm



$t_{i-1}$	$t^0$		$t^1$									$t^2$								
$w_i$	$w^1$	$w^2$	$w^1$			$w^2$			$w^3$			$w^1$			$w^2$			$w^3$		
$w_{i-1}$			$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$
$t_i = t^1$	1.0	1.1	2	4	2	2	2	3	2	4	5	2	2	2	3	2	4	5	2	4
$t_i = t^2$	1.1	2.1	3	5	3	3	2	4	2	4	3	3	3	2	4	2	4	3	4	5

$$= - \sum_k \lambda_k f_k(t_i, t_{i-1}, w_i, w_{i-1})$$


# Viterbi algorithm



$$\phi(t_i, t_{i-1}, w_i, w_{i-1})$$

$t_{i-1}$	$t^0$		$t^1$									$t^2$								
$w_i$	$w^1$	$w^2$	$w^1$			$w^2$			$w^3$			$w^1$			$w^2$			$w^3$		
$w_{i-1}$			$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$
$t_i=t^1$	1.0	1.1	2	4	2	2	2	3	2	4	5	2	2	2	3	2	4	5	2	4
$t_i=t^2$	1.1	2.1	3	5	3	3	2	4	2	4	3	3	3	2	4	2	4	3	4	5

# Viterbi algorithm



$$\phi(t_i, t_{i-1}, w_i, w_{i-1})$$

$t_{i-1}$	$t^0$		$t^1$									$t^2$								
$w_i$	$w^1$	$w^2$	$w^1$			$w^2$			$w^3$			$w^1$			$w^2$			$w^3$		
$w_{i-1}$			$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$	$w^1$	$w^2$	$w^3$
$t_i = t^1$	1.0	1.1	2	4	2	2	2	3	2	4	5	2	2	2	3	2	4	5	2	4
$t_i = t^2$	1.1	2.1	3	5	3	3	2	4	2	4	3	3	3	2	4	2	4	3	4	5



# Training of CRFs: the forward recurrence

---

The conditional log likelihood of the training data under the model is

$$\begin{aligned}\log P_m(T^*|W) &= \sum_{k=1}^N \log \frac{\prod_i e^{\lambda_i f_i(T_k^*, W_k)}}{\sum_{T'} \prod_i e^{\lambda_i f_i(T', W_k)}} \\ &= \sum_{k=1}^N \left( \sum_i \lambda_i f_i(T_k^*, W_k) - \log \sum_{T'} \prod_i e^{\lambda_i f_i(T', W_k)} \right)\end{aligned}$$

We want to compute gradients of this so that we can use gradient descent for optimization

# Training of CRFs: the forward recurrence

$$\sum_{k=1}^N \left( \sum_i \lambda_i f_i (T_k^*, W_k) - \log \sum_{T'} \prod_i e^{\lambda_i f_i (T', W_k)} \right)$$

Sum over all  
training data

Sum of feature potentials;  
easy to compute

Partition function;  
expensive to compute

Use dynamic programming!

# Training of CRFs: the forward recurrence

Define forward variable  $a_i(t^j)$  as the sum of scores of all paths ending up with tag  $t^j$  at word  $w_i$ :

$$a_i(t^j) = \sum_{t' \in \mathcal{T}} \left( e^{\lambda_k f_k(t', t^j, W)} \times a_{i-1}(t') \right)$$

Sum over all possible previous tags

Score for transitioning from previous tag to current tag

Forward variable for previous tag with previous word

# MEMMs and CRFs: Key Points

---

- MEMMs and CRFs incorporate complex features for sequence labeling in a probabilistic framework
- MEMMs and CRFs are based on logistic regression (a.k.a. maximum entropy)
- CRFs improve on MEMMs by using globally, rather than locally normalized probabilities
- CRFs excel at incorporating global constraints on well-formedness of label sequences

---

## [Notebook]

- <https://github.com/scrapinghub/python-crfsuite/blob/master/examples/CoNLL%202002.ipynb>